



Causal inference methods for combining randomized trials and observational studies: a review

Bénédicte Colnet, Imke Mayer, Guanhua Chen, Awa Dieng, Ruohong Li, Gaël Varoquaux, Jean-Philippe Vert, Julie Josse, Shu Yang

► To cite this version:

Bénédicte Colnet, Imke Mayer, Guanhua Chen, Awa Dieng, Ruohong Li, et al.. Causal inference methods for combining randomized trials and observational studies: a review. 2020. hal-03008276

HAL Id: hal-03008276

<https://hal.archives-ouvertes.fr/hal-03008276>

Preprint submitted on 16 Nov 2020

HAL is a multi-disciplinary open access archive for the deposit and dissemination of scientific research documents, whether they are published or not. The documents may come from teaching and research institutions in France or abroad, or from public or private research centers.

L'archive ouverte pluridisciplinaire **HAL**, est destinée au dépôt et à la diffusion de documents scientifiques de niveau recherche, publiés ou non, émanant des établissements d'enseignement et de recherche français ou étrangers, des laboratoires publics ou privés.

Causal inference methods for combining randomized trials and observational studies: a review

Bénédicte Colnet^{1,*}, Imke Mayer^{2,*}, Guanhua Chen³, Awa Dieng⁴, Ruohong Li⁵, Gaël Varoquaux¹, Jean-Philippe Vert⁴, Julie Josse^{6,◊}, Shu Yang^{7,◊}

¹ Université Paris-Saclay, Inria.

² Centre d'Analyse et de Mathématique Sociales, École des Hautes Études en Sciences Sociales, PSL University, CNRS.

³ Department of Biostatistics and Medical Informatics, University of Wisconsin-Madison.

⁴ Google Research, Brain team.

⁵ Department of Biostatistics, Indiana University School of Medicine and Richard M. Fairbanks School of Public Health.

⁶ Inria, Sophia-Antipolis.

⁷ Department of Statistics, North Carolina State University.

* are first co-authors and ◊ are corresponding authors

With increasing data availability, treatment causal effects can be evaluated across different dataset, both randomized trials and observational studies. Randomized trials isolate the effect of the treatment from that of unwanted (confounding) co-occurring effects. But they may be applied to limited populations, and thus lack external validity. On the opposite large observational samples are often more representative of the target population but can conflate confounding effects with the treatment of interest. In this paper, we review the growing literature on methods for causal inference on combined randomized trial and observational studies, striving for the best of both worlds. We first discuss identification and estimation methods that improve generalizability of randomized controlled trials (RCTs) using the representativeness of observational data. Classical estimators include weighting, difference between conditional outcome models, and double robust estimators. We then discuss methods that combining RCTs and observational data to improve the (conditional) average treatment effect estimation, handling possible unmeasured confounding in the observational data. We also connect and contrast works developed in both the potential outcomes framework and the structural causal models framework. Finally, we compare the main methods using a simulation study and real world data to analyse the effect of tranexamic acid on the mortality rate in major trauma patients. Code to implement many of the methods is provided.

Keywords: Causal effect generalization; transportability; external validity; data integration; heterogeneous data; double robustness.

1 Introduction

Experimental data, collected through carefully designed experimental protocols, are usually considered the gold standard approach for assessing the causal effect of an intervention or a treatment on an outcome of interest. Randomized interventions are widely used in many domains such as economics, social sciences and medicine. In particular, the intensive use of randomized controlled trials (RCTs) in the medical area pertains to the so-called “evidence-based medicine”, a keystone of modern medicine. Given that the treatment allocation in an RCT is under control, and that the distribution of covariates for treated and control individuals is often *balanced* for a binary treatment, simple estimators such as the difference in mean effect between the treated and control individuals can consistently estimate the treatment effect (Imbens and Rubin, 2015). However, RCTs can come with drawbacks. First, RCTs can be expensive, take a long time to set up, and be compromised by insufficient sample size due to either recruitment difficulties or restrictive inclusion/exclusion criteria. Second, these criteria for participant eligibility can lead to a narrowly-defined trial sample that differs markedly from the population potentially eligible for the treatment. Therefore, the findings from RCTs can lack generalizability to a target population of interest. This concern is related to the aim of *external validity* central to medical research (Concato et al., 2000; Rothwell, 2005; Green and Glasgow, 2006; Frieden, 2017) and policy research (Martel Garcia and Wantchekon, 2010; Deaton and Cartwright, 2018; Deaton et al., 2019). Note that this concern about the need of generalizability is not shared by all, as discussed by Rothman et al. (2013).

In contrast, there is an abundance of *observational data*, collected without systematically-designed interventions. Such data can come from different sources. For example, they can be collected from research sources such as disease registries, cohorts, biobanks, epidemiological studies, or they can be routinely collected through electronic health records, insurance claims, administrative databases, etc. In that sense, observational data can be readily available, include large samples that are representative of the target populations, and be less costly than RCTs. In order to leverage observational data for causal effect analysis in health domains, the U.S. Food and Drug Administration (FDA) has proposed a framework that distinguishes two types of information levels: Real World Data (RWD) and Real World Evidence (RWE). On the one hand, RWD are data related to individuals’ health status and/or the delivery of health care routinely collected from a variety of sources. On the other hand, RWE is the clinical evidence derived from analysis of RWD. However, there are often concerns about the quality of these “big data” given that the lack of a controlled experimental intervention opens the door to *confounding bias*. In the absence of unknown confounders, there exist many methods to consistently estimate a causal treatment effect from observational data such as matching, inverse propensity weighting (IPW), or augmented IPW (AIPW) (Imbens and Rubin, 2015). However, the *internal validity* of causal claims built from observational data is still not consensual due to the impossibility of completely ruling out confounding bias.

Combining the information gathered from experimental and observational data is a promising avenue to build upon the internal validity of RCTs and a greater external validity of the real-world data. In what follows, we give as examples three potential benefits in real world applications.

Using RCTs and observational data to generalize the treatment effect in a target patient population. The FDA has recently greenlighted the extended usage of a certain drug (Ibrance) to men with breast cancer, though clinical trials performed for authorization on the market were performed on the female population. Breast cancer in men is a rare disease, and

therefore fewer trials were conducted to include male patients. The approval of the extension to the male population was based on the post-market health record and claims data, which cited the real-world usage of Ibrance on men. To authorize such an extension is a solution to reduce the time to approve a drug for patients who could benefit from it. But there are also concerns with such approaches, especially cases where the target population is very different from the RCT cohort. Method development and validation is needed to better navigate the speed-safety balance.

Comparing RCTs and observational data to validate observational methods. There are many research questions for which it is impossible to conduct an RCT, for instance for ethical reasons. Thus, researchers and clinicians remain compelled to take decisions based on observational data. Having at disposal the two sources of data can be useful to benchmark and validate observational studies as there is still concern that the new methods developed to analyze observational data could weaken standards. The Women's Health Initiative (WHI) is a widely-cited example to illustrate potential inconsistencies between conclusions derived from RCTs and RWD. This extensive project was launched to assess whether hormone replacement therapy for healthy women could prevent the appearance of menopausal symptoms. Two out of four interventions of the RCT ended earlier than expected when evidence had accumulated that the health risks exceeded the benefits for this study population (Prentice and Anderson, 2008), whereas observational data first showed a benefit in the hormone replacement therapy. There were many discussions to understand whether the discrepancies between the studies were due to external validity issues of the RCT or internal validity issues of the observational data (Cole and Stuart, 2010; Shadish et al., 2002). This debate led to a reanalysis of the data which showed that the observational data carried the same message as the RCT, and that the differences observed were due to discrepancies in timing of start of treatment and effect over the time (Vandenbroucke, 2009; Frieden, 2017). Following a specific protocol can help avoid such apparent contradictions, by ensuring comparable analyses and a well-stated causal question. More precisely, such a protocol requires first to specify what *target trial* could answer the precise causal question we are interested in, while the observational data is used in a second step to emulate this target trial. The third part of the procedure is to perform sensitivity analysis to investigate the remaining discrepancies (Hernán et al., 2016; Hernán, 2018; Lodi et al., 2019).

Integrating the complementary features of RCTs and observational data to better estimate treatment effects. The COVID-19 health crisis is a timely example of a pandemic where a very rapid response is needed to assess the efficacy of various candidate treatments. In the beginning of the outbreak, there are generally far more observational data than clinical trials. Knowing how to best combine these two sources of information can be crucial, in particular to better estimate heterogeneous treatment effect as RCTs are known to be under-powered in such settings.

There is an abundant literature on the problem of bridging the findings from an experimental study to the target population and combining both sources of information. Similar problems have been termed as *generalizability* (Cole and Stuart, 2010; Stuart et al., 2011; Hernán and VanderWeele, 2011; Tipton, 2013; O'Muircheartaigh and Hedges, 2014; Stuart et al., 2015; Keiding and Louis, 2016; ?; Dahabreh et al., 2019; Buchanan et al., 2018), *representativeness* (Campbell, 1957), *external validity* (Rothwell, 2005; Stuart et al., 2018), *transportability* (Pearl and Bareinboim, 2011; Rudolph and van der Laan, 2017; Westreich et al., 2017), and also *data fusion* (?). They have connections to the covariate shift problem in machine learning (Sugiyama and Kawanabe, 2012). This problem of data integration is tackled in the two main frameworks for causal inference, namely the potential

outcomes (PO) framework (Neyman, 1923; Rubin, 1974), associated with the work by Donald Rubin, and the work on structural causal models (SCM) using directed acyclic graphs (DAGs), much of it associated with work by Judea Pearl (Pearl, 1995) and his collaborators. Note that the DAGs are intuitive tools to clinicians and can be used to incorporate the domain knowledge easily.

The present paper reviews available works on combining experimental data (RCTs) and observational data (RWD). It is organized as follows: In the next section, we introduce the notations in the PO framework, as well as the design setting. The review then starts by considering in Section 3 the case where the available data on the RCT are covariates (also known as baseline covariates when measured at inclusion of the patient), treatment, and outcome, while on observational data, only covariates are available. The aim is to generalize RCTs findings to the target population. To do so, we give the identifiability assumptions and present the main estimation methods, i.e., inverse probability of sampling weighting (IPSW) and stratification, g-formula, doubly robust estimators, etc., that have been suggested to account for distributional shifts. In Section 4, we consider the case where observational data also contain treatment and outcome data. We consider estimators for estimating the conditional average treatment effect using the two data sources while handling potential unmeasured confounders. In Section 5, we present the SCM point of view which can be used to achieve identification when there are many variables. The SCM and PO frameworks are complementary, with different strengths discussed in Imbens (2019) and share many aspects (Richardson and Robins, 2013), e.g, the transportability concept in the SCM framework is equivalent to the generalization (external validity) concept in the PO framework. In Section 6, we present available implementations and software, which we apply using also new implementations, to simulated data. In Section 7, we apply the different methods on a medical application involving major trauma patients where the aim is to assess the effect of the drug tranexamic acid on mortality in head trauma patients and where both an RCT (the CRASH-3 trial) and an observational database (the Traumabase registry) are available. In this section, we also review methods for addressing data quality issues such as missing values.

2 Problem setting

2.1 Notations, in the PO framework

We model each patient in the RCT or observational population as described by a random tuple $(X, Y(0), Y(1), A, S)$ drawn from a distribution \mathcal{P} , where X is a p -dimensional vector of covariates, A denotes the binary treatment assignment (with $A = 0$ for the control and $A = 1$ for the treated patients), $Y(a)$ is the binary or continuous outcome had the subject been given treatment a (for $a \in \{0, 1\}$), and S is a binary indicator for RCT eligibility (i.e., meet the RCT inclusion and exclusion criteria) and willingness to participate if being invited to the trial ($S = 1$ if eligible and also willing to participate if being invited to the trial, $S = 0$ if not)¹.

Assuming *consistency of potential outcomes*, we also denote by $Y = AY(1) + (1 - A)Y(0)$ the outcome realized under treatment A . We model the patients belonging to an RCT sample of size n and in an observational data sample of size m by $n + m$ independent random tuples: $\{X_i, Y_i(0), Y_i(1), A_i, S_i\}_{i=1}^{n+m}$, where the RCT samples $i = 1, \dots, n$ are identically distributed ac-

¹Note that S takes several acceptations depending on research papers, for example other works use two indicators, one for participation and one for eligibility (Nguyen et al., 2018; Dahabreh et al., 2019). In such situations $S = 1$ denotes eligible **and** participating to the trial people, while $S = 0$ denotes ineligible or eligible but necessarily non-participating people

cording to $\mathcal{P}(X, Y(0), Y(1), A, S \mid S = 1)$, and the observational data samples $i = n + 1, \dots, n + m$ are identically distributed according to $\mathcal{P}(X, Y(0), Y(1), A, S)$. For simplicity of exposition, we also denote $\mathcal{R} = \{1, \dots, n\}$ the index set of units observed in the RCT study, and $\mathcal{O} = \{n + 1, \dots, n + m\}$ the index set of units observed in the observational study.

For each RCT sample $i \in \mathcal{R}$, we observe $(X_i, A_i, Y_i, S_i = 1)$, while for observational data $i \in \mathcal{O}$, we consider two settings:

- i. We only observe the covariates X_i ,
- ii. We also observe the treatment and outcome (X_i, A_i, Y_i) .

The former case will be discussed in the next section while the latter will be detailed in Section 4. We define the conditional average treatment effect (CATE):

$$\forall x \in \mathbb{R}^p, \quad \tau(x) = \mathbb{E}[Y(1) - Y(0) \mid X = x],$$

and the population average treatment effect (ATE):

$$\tau = \mathbb{E}[Y(1) - Y(0)] = \mathbb{E}[\tau(X)].$$

We note that in general, if the samples in the RCT and observational data do not follow the same distribution, then the ATE is different from the RCT (or sample) average treatment effect

$$\tau \neq \tau_1, \quad \tau_1 = \mathbb{E}[Y(1) - Y(0) \mid S = 1].$$

We denote respectively by $e(x)$ and $e_1(x)$ the propensity score in the observational data and in the RCT population:

$$e(x) = P(A = 1 \mid X = x), \quad e_1(x) = P(A = 1 \mid X = x, S = 1).$$

We also denote by $\mu_a(x)$ and $\mu_{a,1}(x)$ the conditional mean outcome under treatment $a \in \{0, 1\}$ in the observational data and in the RCT population, respectively:

$$\mu_a(x) = \mathbb{E}[Y(a) \mid X = x], \quad \mu_{a,1}(x) = \mathbb{E}[Y(a) \mid X = x, S = 1],$$

and by $\pi_S(x)$ the selection score²:

$$\pi_S(x) = P(S = 1 \mid X = x).$$

Note that $\pi_S(x)$ is the probability of being *eligible* for selection in the RCT given covariates x . It is different from the probability that an individual with covariates x known to be in the study (RCT or observational population) is selected in the RCT:

$$\pi_S(x) \neq \pi_{\mathcal{R}}(x), \quad \pi_{\mathcal{R}}(x) = P(\exists i \in \mathcal{R}, X_i = x \mid \exists i \in \mathcal{R} \cup \mathcal{O}, X_i = x).$$

We similarly note

$$\pi_{\mathcal{O}}(x) = P(\exists i \in \mathcal{O}, X_i = x \mid \exists i \in \mathcal{R} \cup \mathcal{O}, X_i = x) = 1 - \pi_{\mathcal{R}}(x).$$

Finally, we denote by $\alpha(x)$ the conditional odds ratio that an individual with covariates x is in the RCT or in the observational cohort:

$$\alpha(x) = \frac{P(i \in \mathcal{R} \mid \exists i \in \mathcal{R} \cup \mathcal{O}, X_i = x)}{P(i \in \mathcal{O} \mid \exists i \in \mathcal{R} \cup \mathcal{O}, X_i = x)} = \frac{\pi_{\mathcal{R}}(x)}{\pi_{\mathcal{O}}(x)} = \frac{\pi_{\mathcal{R}}(x)}{1 - \pi_{\mathcal{R}}(x)}.$$

Table 2 summarizes the notations for convenience, and Table 1 illustrates the type of data considered.

²Also named sampling propensity score in Tipton (2013).

Table 1: Illustration of data structure of RCT data (Set \mathcal{R}) and observational data (Set \mathcal{O}) with covariates X , trial eligibility S , binary treatment A and outcome Y . Left: with observed outcomes, Right: with potential outcomes. Note that the S covariate can be either 0 or 1 in the observational data set (while usually unknown from the raw data set in the non-nested design), and is always equals to 1 for observations in the RCT. In the nested design (cf. Section D), S takes only value 0 in the observational data set.

	S	Set	Covariates			Treatment	Outcome	S	Set	Covariates			Treatment	Outcome(s)	
			X_1	X_2	X_3	A	Y			X_1	X_2	X_3	A	$Y(0)$	$Y(1)$
1	1	\mathcal{R}	1.1	20	F	1	1	1	\mathcal{R}	1.1	20	F	1	NA	1
	1	\mathcal{R}	-6	45	F	0	1	1	\mathcal{R}	-6	45	F	0	1	NA
n	1	\mathcal{R}	0	15	M	1	0	1	\mathcal{R}	0	15	M	1	NA	1
$n+1$	0	\mathcal{O}	0	\mathcal{O}
	0	\mathcal{O}	-2	52	M	0	1	0	\mathcal{O}	-2	52	M	0	1	NA
	1	\mathcal{O}	-1	35	M	1	1	1	\mathcal{O}	-1	35	M	1	NA	1
	0	\mathcal{O}	-2	22	M	0	0	0	\mathcal{O}	-2	22	M	0	0	NA

Table 2: List of notations.

Symbol	Description
X	Covariates (also known as baseline covariates when measured at inclusion of the patient)
A	Treatment indicator ($A = 1$ for treatment, $A = 0$ for control)
Y	Outcome of interest
S	Trial eligibility and willingness to participate if invited to ($S = 1$ for eligibility, $S = 0$ for non-eligibility)
n	Size of the RCT study
m	Size of the observational study
\mathcal{R}	Index set of units observed in the RCT study; $\mathcal{R}=\{1,\dots,n\}$
\mathcal{O}	Index set of units observed in the observational study; $\mathcal{O}=\{n+1,\dots,n+m\}$
$\pi_{\mathcal{R}}(x)$	Probability that a unit in $\mathcal{R} \cup \mathcal{O}$ with covariate x is in \mathcal{R}
$\pi_{\mathcal{O}}(x)$	Probability that a unit in $\mathcal{R} \cup \mathcal{O}$ with covariate x is in \mathcal{O}
$\alpha(x)$	Conditional odds ratio $\alpha(x) = \pi_{\mathcal{R}}(x)/\pi_{\mathcal{O}}(x)$
τ	Population average treatment effect (ATE) defined as $\tau = \mathbb{E}[Y(1) - Y(0)]$
τ_1	Trial (or sample) average treatment effect defined as $\tau_1 = \mathbb{E}[Y(1) - Y(0) S = 1]$
$\tau(x)$	Conditional average treatment effect (CATE) defined as $\tau(x) = \mathbb{E}[Y(1) - Y(0) X = x]$
$\tau_1(x)$	Trial conditional average treatment effect defined as $\tau_1(x) = \mathbb{E}[Y(1) - Y(0) X = x, S = 1]$
$e(x)$	Propensity score defined as $e(x) = P(A = 1 X = x)$
$e_1(x)$	Propensity score in the trial defined as $e_1(x) = P(A = 1 X = x, S = 1)$, known by design
$\mu_a(x)$	Outcome mean defined as $\mu_a(x) = \mathbb{E}[Y(a) X = x]$ for $a = 0, 1$
$\mu_{a,1}(x)$	Outcome mean in the trial defined as $\mu_{a,1}(x) = \mathbb{E}[Y(a) X = x, S = 1]$ for $a = 0, 1$
$\pi_S(x)$	Selection score defined as $\pi_S(x) = P(S = 1 X = x)$
$f(X)$	Covariates distribution in the target population
$f(X S = 1)$	Covariates distribution conditional to trial-eligible individuals ($S = 1$)

2.2 Study designs; nested or not

It is important to characterize the study design because the identifiability conditions and therefore the estimators depend on and differ over the designs. Following ? and ?, the study design to obtain the trial and observational samples can be categorized into two types: *nested* study designs and *non-nested* study designs as illustrated on Figure 1. This categorization is similar to the classic one

in survey sampling with multiple datasets. This paper focuses on the **non-nested** design in the main text but we detail identifiability and estimators for the nested case in the Appendix D.

Non-nested trial design involves separate sampling mechanisms for the RCT and the observational samples. The trial sample and the observational sample are obtained separately from the target population(s). Here, we do not limit the two underlying populations for the trial and observational samples to be the same, but we assume that the target populations follow the same superpopulation model. For example, the trial study and the observational study are conducted by different researchers in different times or regions - with small time-specific or regional effects so that the underlying study populations are assumed to follow the same distribution. Note the difference between S and the sets \mathcal{R} and \mathcal{O} , where in the observational sample we can have both $S = 1$ and $S = 0$ (Figure 1). In this review, we consider the case where the observational data set is a random sample *i.i.d* from the target population. The medical application that will be presented in Section 7 corresponds to the non-nested design framework.

Nested trial design involves a two-stage sampling mechanism. First, a large sample is selected from the target population, and then the trial sample is selected from and nested in this sample. The rest of the sample constitutes the observational study. It corresponds to a real medical situation, such as designs for a pragmatic trial embedded in a broader health system. For example, in the Women Health Initiative (WHI), after the end of the initial trial period, a cohort of study participants are followed to measure their long-term outcomes. Olschewski and Scheurlen (1985) introduce the comprehensive cohort study (CCS) design for evaluating competing treatments in which clinically eligible participants are first asked to enroll in a randomized trial and, if they refuse, are then asked to enroll in a parallel observational study in which they can choose treatment according to their own preference, leading to the observational data. Note that in this design, even the causal quantity of interest can be different from the non-nested design, as we may want to transport $\tau_1(x)$ in the observational population such that the causal quantity of interest is $\mathbb{E}[Y(1) - Y(0) | S = 0]$ rather than $\mathbb{E}[Y(1) - Y(0)]$ in the non-nested design.

3 When observational data have no treatment and outcome information

We start by considering the case where only the distribution of the covariates from the observational study is available or used. Note that observational data is a random sample from the target population.

3.1 Assumptions needed to identify the ATE on the target population

A fundamental problem in causal inference is that we can observe at most one of the potential outcomes for an individual subject. In order to identify the ATE from RCT and observational covariate data, we require some of the following assumptions.

3.1.1 Internal validity of the RCT

Assumption 3.1 (Consistency) $Y = AY(1) + (1 - A)Y(0)$.




Figure 1: Schematics of the nested (left) and non-nested (right) designs.

Assumption 3.1 implies that the observed outcome is the potential outcome under the actual assigned treatment.

Assumption 3.2 (Randomization) $Y(a) \perp\!\!\!\perp A | (X, S = 1)$ for all X and $a = 0, 1$.

Assumption 3.2 corresponds to internal validity. It holds by design in a completely randomized experiment, where the treatment is independent of all the potential outcomes and covariates, i.e., $(Y(0), Y(1)) \perp\!\!\!\perp A | S = 1$. It also holds by design in a stratified randomized trial based on a discrete X , where the treatment is independent of all the potential outcomes within each stratum of X : $Y(a) \perp\!\!\!\perp A | (X, S = 1)$. The more general case of conditional randomization is kept.

In addition, in an RCT, it is common that the probability of treatment assignment (also called the propensity score) $e_1(x)$ is known. In a complete randomized trial, the propensity score is fixed as a constant, e.g., $e_1(x) = 0.5$ for all x .

3.1.2 Generalizability of the RCT to the target population

Different assumptions on the generalizability of the RCT to the target population are proposed in the literature, ranging from weak to stringent conditions. We now describe these assumptions and their implications.

Assumption 3.3 (Transportability of the CATE) $\tau_1(x) = \tau(x)$ for all x .

Assumption 3.3 requires that the CATE function is transportable from the RCT to the target population, which is plausible if X captures all the *treatment effect modifiers* and there is no trial

encouragement. It means that the invitation to participate in the trial and trial participation itself do not affect the outcome except through treatment assignment (?).

Assumption 3.4 (Mean exchangeability over treatment assignment and trial participation)

$$\mathbb{E}[Y(a) | X = x, S = 1, A = a] = \mathbb{E}[Y(a) | X = x, S = 1] \text{ (Dahabreh et al., 2019)}$$

$$\text{as well as } \mathbb{E}[Y(a) | X = x, S = 1] = \mathbb{E}[Y(a) | X = x] \text{ for each } x \text{ and } a = 0, 1.$$

Assumption 3.4 means exchangeability over treatment assignment in the trial given covariates is expected to hold by design in the RCT.

Assumption 3.5 (Ignorability assumption on trial participation) $\{Y(0), Y(1)\} \perp\!\!\!\perp S | X$.
(Stuart et al., 2011; Buchanan et al., 2018)

Assumption 3.5 means that being eligible to the RCT does not affect the potential outcomes conditionally on the covariates X . In other words, it means that there are no unmeasured variables associated with the trial sample selection and the treatment effect. This is the parallel with the *ignorability assumption* on treatment assignment in causal inference with observational data (see 4.1), but with the sample selection.

It is worth discussing the relationships among Assumptions 3.3 – 3.5, ranging from weak to strong conditions. Assumption 3.3, *transportability of the CATE*, is implied by, but does not imply, the stronger conditions in Assumption 3.4. Assumption 3.4, *mean exchangeability over treatment assignment*, is implied by, but does not imply, the stronger conditions $Y(a) \perp\!\!\!\perp A | X, S = 1$ implied by Assumption 3.2. In the same way, *mean exchangeability over trial participation*, also known as *mean generalizability* (from trial participants to the target population), is implied by, but does not imply, the stronger condition $Y(a) \perp\!\!\!\perp S | X$ in Assumption 3.5 (participation in the RCT is randomized within levels of X). Note that for binary S , mean exchangeability over trial participation implies the mean transportability condition $\mathbb{E}[Y(a) | X = x, S = 1] = \mathbb{E}[Y(a) | X = x, S = 0]$ in Assumption 3.3, for all x such that $P(S = 0 | X = x) > 0$.

Assumption 3.6 (Positivity of trial participation) There exists a constant c such that for all x with probability 1, $\pi_S(x) \geq c > 0$; and $0 < P(A = a | X = x, S = 1) < 1$ for all a and for all x such that $P(S = 1 | X = x) > 0$.

Assumption 3.6 means that we require adequate overlap of the covariate distribution between the trial sample and the target population (in other words, all members of the target population have nonzero probability of being selected into the trial), and also between the treatment groups over the trial sample. The positivity of treatment assignment in the trial given covariates, related to the assumption required for causal inference in confounded settings, is expected to hold by design in the RCT.

3.1.3 Towards identification formula

Under Assumptions 3.1, 3.4 (or 3.5), and 3.6 (for the regression formulation), and under Assumptions 3.1, 3.3 (or 3.4 or 3.5), and 3.6 (for the reweighting formulation), the ATE can be identified based on the following formula (proved in Appendix C):

- a) Reweighting formulation:

$$\tau = \mathbb{E} \left[\frac{n}{m\alpha(X)} \tau_1(X) | S = 1 \right] = \mathbb{E} \left[\frac{n}{m\alpha(X)} \left\{ \frac{A}{e_1(X)} - \frac{1 - A}{1 - e_1(X)} \right\} Y | S = 1 \right]. \quad (1)$$

In a fully randomized RCT where $e_1(x) = 0.5$ for all x , this formula further simplifies to

$$\tau = \mathbb{E} \left[\frac{2n}{m\alpha(X)} (2A - 1)Y \mid S = 1 \right].$$

b) Regression formulation:

$$\tau = \mathbb{E} [\mu_{1,1}(X) - \mu_{0,1}(X)]. \quad (2)$$

Different identification formulas motivate different estimating strategies as discussed in the next subsection.

3.2 Estimation methods to improve generalizability of RCT analysis

As the RCT assigns treatment at random to the participants, the CATE $\tau_1(x) = \tau(x)$ is identifiable (under assumptions 3.1 and 3.2) and can be estimated by standard estimators, such as the difference in means solely from the RCT (see Appendix A). However, in general, the covariate distribution of the RCT sample $f(X \mid S = 1)$ is different from that of the target population $f(X)$; therefore, τ_1 is different from τ in general, and a SATE estimator using only trial data is biased for the ATE of interest.

3.2.1 IPSW and stratification: modeling the probability of trial participation

To overcome this bias, most existing methods rely on direct modeling of the selection score previously introduced. The selection score adjustments methods include inverse probability of sampling weighting (IPSW; Cole and Stuart, 2010; Stuart et al., 2011; Buchanan et al., 2018) and stratification (Stuart et al., 2011; Tipton, 2013; O’Muircheartaigh and Hedges, 2014).

Inverse probability of sampling weighting (IPSW) The IPSW approach can be seen as the counterpart of inverse propensity weighting (IPW) methods to estimate the ATE from observational studies to control for confounding (see Appendix B for details). Based on the identification formula (Equation 1), the IPSW estimator of the ATE is defined as the weighted difference of average outcomes between the treated and control group in the trial. The observations are weighted by the inverse odds ratio $1/\alpha(x) = \pi_O(x)/\pi_R(x)$ to account for the shift of the covariate distribution from the RCT sample to the target population. The larger π_R , the smaller the weight of the observation. The shape of the IPSW estimator is slightly different from the shape of the IPW estimator. In the latter, each observation is weighted by the inverse of the probability to be treated whereas in the former it is by the inverse of the odds ratio of the probability to be in the trial sample. This is due to the non-nested sampling design (see the estimator for the nested design (S3)), as mentioned by Kern et al. (2016) and Nguyen et al. (2018). The IPSW estimator can be written as follows:

$$\hat{\tau}^{\text{IPSW}} = \frac{1}{m} \sum_{i=1}^n \frac{Y_i}{\hat{\alpha}(X_i)} \left(\frac{A_i}{e_1(X_i)} - \frac{1 - A_i}{1 - e_1(X_i)} \right), \quad (3)$$

where $\hat{\alpha}$ is an estimate of α . In a standard RCT with $e_1(x) = 0.5$ being fixed, this further simplifies to

$$\hat{\tau}^{\text{IPSW}} = \frac{1}{m} \sum_{i=1}^n \frac{2Y_i(2A_i - 1)}{\hat{\alpha}(X_i)}.$$

The IPSW estimator is consistent when the quantity α is consistently estimated by $\hat{\alpha}$. The most common method for estimating the selection score is to define it as $\hat{\alpha}(x) = \hat{\pi}_{\mathcal{R}}(x)/(1 - \hat{\pi}_{\mathcal{R}}(x))$, where $\hat{\pi}_{\mathcal{R}}$ is estimated by logistic regression trained to discriminate RCT from observational samples (Stuart, 2010), while recent works also use other methods such as random forest and GBM (Kern et al., 2016). Similar to IPW estimators, IPSW estimators are known to be highly unstable, especially when the selection scores are extreme. It happens when the trial study contains units with very small probabilities of being in the trial. Normalized weights can be used to overcome this issue (?). Still, the major challenge is that IPSW estimators require a correct model specification of the selection score. Avoiding this problem requires either very strong domain expertise or turning to doubly robust methods (Section 3.2.4). Finally, Dahabreh et al. (2019) propose the use of sandwich-type variance estimators (for both nested and non-nested design) or non-parametric bootstrap approaches, and note that the latter may be preferred in practice.

Stratification The stratification approach – also called subclassification – is introduced by Cochran (1968), and further detailed by Stuart et al. (2011), Tipton (2013), and O’Muircheartaigh and Hedges (2014). It is proposed as a solution to mitigate the dangers of extreme weights in the IPSW formula. The principle is to form L strata based on the covariate values via the selection score, which in practice consists in grouping observations for which $\hat{\pi}_{\mathcal{R}}$ are similar. Then, in each stratum l , the average treatment effect is estimated as $\bar{Y}(1)_l - \bar{Y}(0)_l$ where $\bar{Y}(a)_l$ denotes the average value of the outcome for units with treatment a in stratum l in the RCT. Finally, the ATE estimator is the weighted sum of the difference between means in each stratum:

$$\hat{\tau}^{\text{strat}} = \sum_{l=1}^L \frac{m_l}{m} \left\{ \bar{Y}(1)_l - \bar{Y}(0)_l \right\}, \quad (4)$$

where m_l/m is the proportion of population cases in the stratum l in the observational study. Kang et al. (2007) assess performance of the stratification as opposed to the IPSW, and show that stratification is less effective than IPSW at removing bias, but that IPSW struggle more than stratification when the selection scores are small.

3.2.2 G-formula estimators: modeling the conditional outcome in the trial

An alternative to IPSW for generalizing RCT findings to a target population consists in leveraging the regression formulation (Equation 2). The corresponding estimators, known as g -formula estimators, fit a model of the conditional outcome mean among trial participants, instead of fitting a model for the probability of trial membership. Applying these models to the covariates in the target population, i.e., marginalizing over the empirical covariate distribution of the target population, gives the corresponding expected outcome (Robins, 1986). This outcome model-based estimator is then defined as:

$$\hat{\tau}^g = \frac{1}{m} \sum_{i=n+1}^{n+m} (\hat{\mu}_{1,1}(X_i) - \hat{\mu}_{0,1}(X_i)), \quad (5)$$

where $\hat{\mu}_{a,1}(X_i)$ is an estimator of $\mu_{a,1}(X_i)$. In the simplest case, one can assume a linear regression model for each treatment a , estimate it by standard ordinary least squares (OLS), on the trial sample and apply it on the observational sample. When the model is correctly specified the estimator is consistent. Kern et al. (2016) suggest using BART to learn the regression functions.

Dahabreh et al. (2019) show that g -estimators are equivalent to specific IP weighting estimators where the probability of treatment among trial participants is also modeled and when all models are estimated by nonparametric frequency (non-smooth) estimators. Nevertheless, their method encounters difficulties in practice, especially in high dimension, so that they suggest resorting to double robust estimators (Section 3.2.4).

3.2.3 Calibration weighting: balancing covariates

Beyond propensity scores, other schemes use sample reweighting. Dong et al. (2020) propose a calibration weighting approach, which is similar to the idea of entropy balancing weights introduced by Hainmueller (2012). They calibrate subjects in the RCT samples in such a way that after calibration, the covariate distribution of the RCT sample empirically matches the target population. Let $\mathbf{g}(X)$ be a vector of functions of X to be calibrated; e.g., the moments, interactions, and non-linear transformations of components of X . In order to calibrate, they assign a weight ω_i to each subject i in the RCT sample by solving the optimization problem:

$$\min_{\omega_1, \dots, \omega_n} \sum_{i=1}^n \omega_i \log \omega_i, \quad (6)$$

subject to $\omega_i \geq 0$, for all i ,

$$\sum_{i=1}^n \omega_i = 1, \sum_{i=1}^n \omega_i \mathbf{g}(X_i) = \tilde{\mathbf{g}}, \text{ (the balancing constraint)}$$

where $\tilde{\mathbf{g}} = m^{-1} \sum_{i=n+1}^{m+n} \mathbf{g}(X_i)$ is a consistent estimator of $\mathbb{E}[\mathbf{g}(X)]$ from the observational sample. The balancing constraint calibrates the covariate distribution of the RCT sample to the target population in terms of $\mathbf{g}(X)$. The objective function in (6) is the negative entropy of the calibration weights; thus, minimizing this criteria ensures that the empirical distribution of calibration weights are not too far away from the uniform, such that it minimizes the variability due to heterogeneous weights. This optimization problem can be solved using convex optimization with Lagrange multipliers. For an intuitive understanding of the calibration weighting framework, consider $\mathbf{g}(X) = X$. In such a setting, the balancing constraint is forcing the means of the observational data and RCT to be equal after reweighting. More complex constraints can for instance balance higher-order moments. The calibration algorithm is inherently imposing a log-linear model of the sampling propensity score and solving the corresponding parameters by a set of estimating equations induced from covariate balance. This equivalence provides insights of using the penalized estimating equation approach to select important variables for balancing.

Based on the calibration weights, the CW estimator is then

$$\hat{\tau}^{\text{CW}} = \sum_{i=1}^n \hat{\omega}_i Y_i \left\{ \frac{A_i}{e_1(X_i)} - \frac{1 - A_i}{1 - e_1(X_i)} \right\}. \quad (7)$$

The CW estimator $\hat{\tau}^{\text{CW}}$ is doubly robust in that it is a consistent estimator for τ if the selection score of RCT participation follows a loglinear model, i.e., $\pi_S(X) = \exp\{\boldsymbol{\eta}_0^\top \mathbf{g}(X)\}$ for some $\boldsymbol{\eta}_0$, or if the CATE is linear in $\mathbf{g}(X)$, i.e., $\tau(X) = \boldsymbol{\gamma}_0^\top \mathbf{g}(X)$, though not necessarily both. The authors suggest a bootstrap approach to estimate its variance.

3.2.4 Doubly-robust estimators

The model for the expectation of the outcomes among randomized individuals (used in the g -estimators in Equation (5)) and the model for the probability of trial participation (used in IPSW estimators in Equation (3)) can be combined to form an Augmented IPSW estimator (AIPSW):

$$\begin{aligned}\hat{\tau}^{\text{AIPSW}} = & \frac{1}{m} \sum_{i=1}^n \frac{1}{\hat{\alpha}(X_i)} \left[\frac{A_i \{Y_i - \hat{\mu}_{1,1}(X_i)\}}{e_1(X_i)} - \frac{(1 - A_i) \{Y_i - \hat{\mu}_{0,1}(X_i)\}}{1 - e_1(X_i)} \right] \\ & + \frac{1}{m} \sum_{i=n+1}^{m+n} \{\hat{\mu}_{1,1}(X_i) - \hat{\mu}_{0,1}(X_i)\}.\end{aligned}$$

It can be shown to be doubly robust, i.e., consistent and asymptotically normal when either one of the two models for $\hat{\pi}_R$ and $\hat{\mu}_{a,1}(X)$ ($a = 0, 1$) is correctly specified, as demonstrated by ?, supplementary material for both the nested and non-nested designs.

More recently, Dong et al. (2020) propose an augmented calibration weighting (ACW) estimator, given by

$$\begin{aligned}\hat{\tau}^{\text{ACW}} = & \sum_{i=1}^n \hat{\omega}_i \left[\frac{A_i \{Y_i - \hat{\mu}_{1,1}(X_i)\}}{e_1(X_i)} - \frac{(1 - A_i) \{Y_i - \hat{\mu}_{0,1}(X_i)\}}{1 - e_1(X_i)} \right] \\ & + \frac{1}{m} \sum_{i=n+1}^{m+n} \{\hat{\mu}_{1,1}(X_i) - \hat{\mu}_{0,1}(X_i)\}.\end{aligned}\tag{8}$$

They show that $\hat{\tau}^{\text{ACW}}$ achieves double robustness and local efficiency, i.e., its asymptotic variance achieves the semiparametric efficiency bound (the variance is smaller than the asymptotic variance of $\hat{\tau}^{\text{CW}}$ in equation 7). Moreover, the ACW estimator enables the use of double machine-learning estimation of nuisance functions (estimates of both) while preserving the root-n consistency of the ACW estimator. Other doubly robust estimators include targeted maximum likelihood estimators (TMLE) (Rudolph and van der Laan, 2017). They provide a semiparametric efficiency score for transporting the ATE from one study site to another, where one site is regarded as a population. ? propose double/debiased machine learning methods to consistently estimate the ATE by using flexible machine learning methods for the nuisance parameters estimation to avoid model mis-specification. Their approach uses Neyman-orthogonal scores to reduce sensitivity of ATE estimation with respect to nuisance parameters. In addition, they use cross-fitting (Zheng and van der Laan, 2011; ?) to provide an efficient and unbiased form of data-splitting.

Dahabreh et al. (2020) perform a simulation study when all hypothesis are met to compare IPSW, g -estimators and doubly robust ones. It confirms that when all models are correctly specified all estimators are approximately unbiased, with the outcome-model based estimator (5) showing the lowest variance. Note that they do not explicitly simulate under model's mis-specification.

We must add the caveat that all methods assume the transportability condition: given the covariates x the treatment effect must be the same in the observational data as in the trial. It could be broken if some treatment effect modifiers are not captured in the data. However, if the observational data provide additional treatment and outcome information, some key assumptions may be testable. In the next section, we review methods developed in the context of combining RCT and full observational data.

4 When observational data contain treatment and outcome information

In Section 3, we have studied the use of the covariate distribution of the observational sample to adjust for selection bias of the RCT sample. Now we consider the setting where we have access to additional treatment and outcome information (Y, A) from the observational sample. Many studies involve both RCT and observational data with comparable information, e.g., the study we detail in Section 7 with the Traumabase and the CRASH-3 datasets. In this context, the question of interest becomes how to leverage both data sources for efficient estimation of the ATE and CATE.

4.1 Causal inference on observational data

Under classical identifiability assumptions, it is possible to estimate the ATE and CATE based only on the observational data. The classical ones are the following.

Assumption 4.1 (Unconfoundedness) $Y(a) \perp\!\!\!\perp A | X$ for $a = 0, 1$.

Assumption 4.1 (also called *ignorability* assumption) states that treatment assignment is as good as random conditionally on the attributes X . In other words, all confounding factors are measured. Unlike the RCT, in observational studies, its plausibility relies on whether or not the observed covariates X include all the confounders that affect the treatment as well as the outcome.

Assumption 4.2 (Overlap) There exists a constant $\eta > 0$ such that for almost all x , $\eta < e(x) < 1 - \eta$.

Assumption 4.2 (also called *positivity* assumption) states that the propensity score $e(\cdot)$ is bounded away from 0 and 1 almost surely.

Under Assumptions 4.1 and 4.2, the ATE can be identified based on the following formula from the observational data:

a) Reweighting formulation:

$$\tau = \mathbb{E} \left[\frac{AY}{e(X)} - \frac{(1-A)Y}{1-e(X)} \right]; \quad (9)$$

b) Regression formulation:

$$\tau = \mathbb{E} [\tau(X)] = \mathbb{E} [\mu_1(X) - \mu_0(X)]. \quad (10)$$

The identification formulas motivate IPW (Appendix B), regression estimators or doubly robust estimators based solely on the observational data. There are many methods also available to estimate the CATE $\tau(\cdot)$ based on the observational data such as causal forests (Wager and Athey, 2018), causal Bayesian additive trees (BART) (Hill, 2011; Hahn et al., 2020), causal boosting (?), or causal multivariate adaptive regression splines (MARS) (?). There are also meta-learners such as the S-Learner (Künzel et al., 2018), T-learner (Künzel et al., 2018), X-Learner (Künzel et al., 2019), MO-Learner (Rubin and van der Laan, 2007; Künzel et al., 2018), and R-learner (Nie and Wager, 2017), which build upon any base learners for regression or supervised classification such as random forests (Breiman, 2001) and BART (Hill, 2011).

4.2 Dealing with unmeasured confounders in observational data

The ATE is useful to inform about the treatment effect over a target population on average. Alternatively, the CATE informs how treatment effect varies over individual characteristics. Under Assumptions 3.5 and 4.1, the CATE can be estimated based on the RCT and observational study, and therefore the two data sources can be pooled to improve estimation efficiency. In this case, the covariate shift problem for the RCT does not bias the estimator of the CATE; furthermore, we do not require the covariate distribution to overlap between the RCT and observational data. In practice, observational data may violate the desirable assumption for combining, e.g., that there are no unmeasured confounders. The design benefit of RCTs can be used to overcome this lack of internal validity. Hence, the next questions are whether we should include observational data into our analysis and how we should use it. At a general level, we face a case where we want to combine an unbiased but high-variance estimator (due to the small sample size of the RCT) and a biased but low-variance estimator from the observational study.

To answer the first question, Yang et al. (2020) derive a statistical test to gauge the reliability of the observational data compared to the gold-standard RCT data. The test outcome determines whether or not to use the observational data in an integrative analysis. Their strategy leads to an elastic test-based integrative estimator that uses the optimal combining rule for estimation if the violation test is not significant and retains only the RCT counterpart if the violation test is significant. This guarantees the consistency of the CATE estimator regardless of whether or not the observational data meet the criteria for combining. The elastic integrative estimator gains efficiency over the RCT alone estimator and gains robustness to unmeasured confounding over the naive combining estimator.

Other approaches exist to handle unmeasured confounders. Kallus et al. (2018) consider a setting where the ignorability on the trial assumption (Assumption 3.5) holds but where the observational data does not fully overlap with the RCT. They suggest the following estimation strategy. First, using confounded observational data $\{(X_j, A_j, Y_j) : j \in \mathcal{O}\}$, they estimate the conditional treatment effect with classical methods such as causal forest (Wager and Athey, 2018), denoted by $\hat{\omega}(x)$. Due to possible unmeasured confounding, $\hat{\omega}(x)$ may be biased for $\tau(x)$. To correct for this bias, they write the bias as $\eta(x) = \tau(x) - \hat{\omega}(x)$. Given that $\hat{\omega}(x)$ is obtained from the observational data, one can learn the bias term η using unconfounded RCT data, $\{(X_i, A_i, Y_i, S_i = 1) : i \in \mathcal{R}\}$. Furthermore, they assume that the bias can be well approximated by a function with low complexities: e.g., a linear function of x : $\eta(x) = \theta^T x$. This assumption guarantees the validity of the framework even if the observational data does not fully overlap with the experimental data as the bias can be well estimated by extrapolating. Lack of overlap could be an issue if the bias is approximated using nonparametric models such as random forests because their ability to extrapolate may be weak. Kallus et al. (2018) then use the transformed outcome $Y_i^* = [e(X_i)^{-1} A_i - \{1 - e(X_i)\}^{-1} (1 - A_i)] Y_i$, which is unbiased for $\tau(X_i)$, $\mathbb{E}[Y_i^* | X_i] = \tau(X_i)$ and estimate the bias as $\hat{\eta}(x) = \hat{\theta}^T x$ where:

$$\hat{\theta} = \operatorname{argmin}_{\theta} \sum_{i=1}^n \{Y_i^* - \hat{\omega}(X_i) - \theta^T X_i\}^2.$$

Finally, $\hat{\tau}(x) = \hat{\omega}(x) + \hat{\eta}(x)$ is the estimated conditional average treatment effect. They prove that under conditions of parametric identification of η , they recover a consistent estimate of $\tau(x)$ at a rate which is governed by the rate of estimating ω by $\hat{\omega}$.

In clinical settings, parametric models for the CATE are desirable due to their easy interpretations. Yang et al. (2020) consider a structural model assumption for the CATE, e.g., for continuous

outcomes, a linear CATE function of the form $\tau_{\varphi_0}(x) = x^T \varphi_0$ with $\varphi \in \mathbb{R}^{p_1}$, and for binary outcomes, a CATE function of the form $\tau_{\varphi_0}(x) = \{\exp(x^T \varphi_0) - 1\}/\{\exp(x^T \varphi_0) + 1\}$ ranging from -1 to 1 . Furthermore, to accommodate the possible unmeasured confounders in the observational data, they assume that a confounding function summarizes the impact of unmeasured confounders on the difference in the potential outcome between the treated and untreated groups, accounting for the observed covariates. In particular, they consider

$$\lambda(x) = \mathbb{E}[Y(0) | A = 1, X = x] - \mathbb{E}[Y(0) | A = 0, X = x]$$

for the observational data. In the absence of unmeasured confounding $\lambda(x) = 0$ for any $x \in \mathbb{R}^{p_1}$, while if there is unmeasured confounding, $\lambda(x) \neq 0$ for some x . Similar to the CATE, Yang et al. (2020) impose a structural model assumption for $\lambda(x) = \lambda_{\phi_0}(x)$ with $\phi \in \mathbb{R}^{p_2}$. The confounding function is unidentifiable based only on the observational data. Coupling the RCT and observational data, Yang et al. (2020) show that the CATE and confounding function are identifiable. The key insight is to define the random variable

$$H_{\psi_0} = Y - \tau_{\varphi_0}(X)A - (1 - S)\lambda_{\phi_0}(X)\{A - e(X)\}, \quad (11)$$

where $\psi_0 = (\varphi_0^T, \phi_0^T)^T \in \mathbb{R}^p$ ($p = p_1 + p_2$) is the full vector of model parameters in the CATE and confounding function. By separating the treatment effect $\tau_{\varphi_0}(X)A$ and $(1 - S)\lambda_{\phi_0}(X)\{A - e(X)\}$ from the observed Y , H_{ψ_0} mimics the potential outcome $Y(0)$. They then derive the semiparametric efficient score of ψ_0 :

$$S_{\psi_0}(V) = \begin{pmatrix} \frac{\partial \tau_{\varphi_0}(X)}{\partial \varphi_0} \\ \frac{\partial \lambda_{\phi_0}(X)}{\partial \phi_0}(1 - S) \end{pmatrix} \{\sigma_S^2(X)\}^{-1} (H_{\psi_0} - \mathbb{E}[H_{\psi_0} | X, S])\{A - e(X)\}, \quad (12)$$

where $\sigma_S^2(X) = \mathbb{V}[Y(0) | X, S]$. A semiparametric efficient estimator of ψ_0 can be obtained by solving the estimating equation based on (12). If the predictors in $\tau_{\varphi_0}(X)$ and $\lambda_{\phi_0}(X)$ are not linearly dependent, they show that the integrative estimator of the CATE is strictly more efficient than the RCT estimator. As a by-product, this framework can be used to generalize the ATEs from the RCT to a target population without requiring an overlap covariate distribution assumption between the RCT and observational data.

Methods mentioned above are applicable for general cases of integrative analysis, other approaches have been tailored for special applications. Peysakhovich and Lada (2016) propose a method for integrative analysis when observational data are time series. First, one can use observational time series data to estimate a mapping from observed treatments to unit-level effects. This estimate is biased due to potential unobserved confounders. Then, one can use experimental data to identify a monotonic transformation from biased estimates to real treatment effects. To use this method, unit-level time series data are needed for the first step and assume the bias preserves unit-level relative rank ordering. Athey et al. (2020) combine RCT and observational data to obtain credible estimates of the causal effect on a primary outcome in a setting where both observational and RCT samples contain treatment, features, and a secondary (often short-term) outcome, but the primary outcome is observed only in the non-randomized sample, the rational being that the treatment effect on the secondary outcome and that on the primary outcome should be similar. If this is not the case, they assume that it is because of unobserved confounders in the observational sample. Their method consists in adjusting the estimates of the treatment effects on the primary outcome using the differences observed on the secondary outcome. They suggest three methods,

namely, *i*) imputing the missing primary outcome in the RCT, *ii*) weighting the units in the observational sample to remove biases and *iii*) using control function methods. The key assumption for identifiability is a latent unconfoundedness $Y_i^P(a) \perp\!\!\!\perp W_i \mid Y_i^S(a), X_i, i \in \mathcal{O}$, i.e., unobserved confounders that affect both treatment assignment and the secondary outcome in the observational study are the same unobserved confounders that affect both the treatment assignment and primary outcome. Their assumptions also imply that the missing data in the potential outcomes are missing at random (Rubin, 1976).

5 Structural causal models

The PO framework and the associated methods reviewed in the previous sections rely on the implicit assumption that we know which variable – the treatment assignment – is susceptible to causally affect another variable – the outcome – and which variables could confound this effect – the confounders. It also assumes that we know which variables are responsible for trial allocation. Consequently, the PO framework deals with identifiability of the resulting treatment effect by postulating assumptions about the data generating process (ignorability, overlap, treatment consistency, etc.) and extensively discusses its estimation. In certain applications, systems are too complex or involve too many covariates to assume a complete *a priori* knowledge about these different components. The Structural Causal Models (SCM) framework offers an alternative for identifiability assessment by proposing a different formalization of the problem. This formalization allows to identify sets of relevant variables to rule out different types of biases, while still incorporating all available domain knowledge, and to establish identifiability of the treatment effect as we will see in the following.

In this section, we review the SCM framework for combining RCTs and observational data, focusing particularly on the recent work by ? and Hünermund and Bareinboim (2019) as they summarize this line of work (see reference therein). The authors propose a unified framework called *data fusion* which accounts for the heterogeneity of data, i.e., the fact that “not all data are created equally”. More precisely, they propose a general framework, illustrated in Figure 2, that simultaneously takes into account confounding issues of observational data, sample selection issues, as well as extrapolation of causal claims across heterogeneous settings, in a *non-parametric* setting. Most of their work focuses on identifiability issues to determine whether a query Q , often the treatment effect, can be estimated from the available information. Links with the PO framework are underlined. Finally, we discuss the connection and complementarity between the SCM framework and the PO framework. Of note, the exposition here relies on so-called structural equation models with independent errors and their representation using directed acyclic graphs (DAGs), including DAGs with selection nodes (?); an alternative approach for generalizability and integrative analyses of trials and observational studies (Dahabreh et al., 2019, 2020) using structural equation models under weaker error assumptions and represented using single world intervention graphs (Richardson and Robins, 2013) has also been proposed but will not be examined in detail here.

5.1 Data fusion: a unified framework for combining observational data and RCT

Note that we adapt the notations of the SCM framework to fit with those introduced in previous sections. Specifically, we denote Y an outcome, A a treatment variable, X observed covariates, U unobserved variables and S a sample selection indicator. We provide a short introduction to SCM and *do-calculus* (Pearl, 2009) in Appendix E.

Given a target quantity to estimate (mostly $P(y | do(a))$), ? and Hünermund and Bareinboim (2019) introduce a generalized framework which identifies when the target quantity is estimable using available data including for instance purely observational data (giving $P(a, x, y)$), other interventional data, and nonexperimental studies (giving $P(a, x, y, | S = 1)$). The framework relies on defining assumptions and conditions under which the iterative application of *do-calculus* rules can lead to the reduction of the target quantity onto an estimable expression (containing only standard probabilities, i.e., *do-less* expressions).

Figure 2 gives a schematic illustration of the proposed framework. In what follows, we summarize the conditions and identifiability results proposed by the authors for the cases of handling confounding bias, selection bias, and generalizability of RCT results to a target population.




Figure 2: Schematic illustration of the data fusion process. The causal inference engine provided by *do-calculus* takes three inputs: (1) a causal effect query Q , (2) a model G , and (3) the type of data, $P(v | \cdot)$, that is available. It is guaranteed to return a transformation of Q , based on G , that is estimable with the available data, whenever such a solution exists.

5.1.1 Confounding bias

In order to estimate the causal effect $P(y | do(a))$ using only available observational data, following the observational distribution $P(a, x, y)$, the idea is to identify—on the basis of the causal graph—a set of *admissible variables* such that measuring and adjusting for these variables removes any bias due to confounding. The *backdoor criterion* defined below provides a graphical method for selecting admissible sets for adjustment.

Definition 1 (Admissible sets - the backdoor criterion) Given an ordered pair of treatment and outcome variables (A, Y) in a causal DAG G , a set X is backdoor admissible if it blocks every path between A and Y in the graph G_A , with G_A the graph that is obtained when all edges emitted by node A are deleted in G .

The backdoor criterion can be seen as the counterpart of unconfoundedness in Assumption 4.1: If a set X of variables satisfies the backdoor condition relative to (A, Y) , then $Y(a) \perp\!\!\!\perp A | X$. Identifying backdoor admissible variables is important because it allows to estimate causal effects from observational data as follows:

Theorem 5.1 (Backdoor adjustment criterion) If a set of variables satisfies the backdoor criterion relative to (A, Y) , the causal effect of A on Y can be identified from observational data by the adjustment formula $P(y | do(a)) = \sum_x P(y | a, x)P(x)$.

The adjustment formula can be seen as the counterpart of the identifiability formula in Equation 10.

The backdoor criterion is one of the graphical methods for identifying admissible sets. In cases where it is not applicable, an extended definition called the *frontdoor criterion* can be applied using mediators in the graph. Figure 3 provides a summary of the identifiability conditions when the available data is either observational data or data from surrogate experiments.





Figure 3: **Summary of identifiability results to control for confounding bias:** If there exists a set of observed variables that satisfies the backdoor criterion, then the causal effect of A on Y can be identified using nonexperimental data alone. In the case where no set of observed variables satisfies the backdoor condition but the effect of A can be mediated by an observed variable M (mediator), if there exists a set of observed variables that satisfies the frontdoor criterion, then the causal effect is also identifiable from observational data alone. If none of these conditions holds, the query is not identifiable. If, in addition to observational data, RCTs through surrogate experiments are available, the z -identifiability condition is sufficient to determine if the query is identifiable or not.

5.1.2 Sample selection bias

To tackle sample selection bias, i.e., preferential selection of units, the authors consider an indicator variable S such that $S = 1$ identifies units in the sample. The data at hand can be seen as $P(a, y, x \mid S = 1)$ and the target is $P(y \mid do(a))$. Figure 4 (b) presents a case where the selection



process is d -separated (definition in Appendix E) from Y by A , then $P(y \mid a) = P(y \mid a, S = 1)$; since A and Y are unconfounded, $P(y \mid do(a)) = P(y \mid a)$ so that the experimental distribution is recoverable from observed data. This is not the case for Figure 4 (a) without further assumptions. When both confounding bias and selection bias are present in the data (Figure 4 (c)), the graphical framework can help selecting among the list of adjustment sets, $\{W_1, W_2\}$, $\{W_1, W_2, X\}$, $\{W_1, X\}$, $\{W_2, X\}$, and X , (these sets control for confounding), the one that can be used as available from biased data; here it will be X as it is the only one separated from S , leading to $P(y \mid do(a)) = \sum_x P(y \mid a, x, S = 1)P(x \mid S = 1)$. This ability to select relevant covariates for identifiability is presented as an important advantage of the SCM framework.

Combined bias and unbiased data. Note that the previous examples in Figure 4 concern only one set of data but the approach is extended to combine data, biased (with a selection) data, and unbiased data (for instance covariates from the target population) as follows. To do so, ? define the **S -backdoor admissible criterion** which is a sufficient condition but not necessary. It states that if X is backdoor admissible, A and X block all paths between S and Y , i.e. $Y \perp\!\!\!\perp S \mid A, X$, and that X is measured in both population-level data and biased data, then, the causal effect can be identified as

$$P(y \mid do(a)) = \sum_x P(y \mid a, x, S = 1)P(x).$$

This expression shows that one can generalize what is observed on the selected sample by reweighting or recalibrating by $P(x)$ that is available from the target population (unbiased data). More complex setting can be handled, such as dealing with post-treatment variables. In such a case, they show that generalizability can be obtained by another weighting strategy (not by $P(x)$), which can also be seen as a benefit of this framework.

5.1.3 Transportability: extrapolating causal knowledge across domain

Finally, ? tackle the issue of transportability, which consists in extrapolating causal knowledge across domain while settings' populations and environments differ both in their distributions and in their inherent causal characteristics. For example, transportability could help answering questions such as: “If a program worked for poor rural women in Africa, will it work for middle-income urban

men in South Asia?”. This range of questions address the external validity issue presented in the introduction.

? introduce a formal representation called “selection diagrams” to express the differences and commonalities between two samples as represented in Figure 5. Causal diagrams are augmented with a set S of “selection variables”, where each member of S corresponds to a mechanism by which the two samples are different (either in the distribution of background factors (U) or due to divergent causal mechanisms f defined in Appendix E). Note that there is a square for S instead of a nested circle for selection bias like in Figure 4. Note that S is emitting arrows, whereas selection nodes indicating preferential inclusion into the sample only receive incoming arrows.

Let us consider two domains, $S = s^*$ and $S = s$, in the graph of Figure 5 (a) and focus on the aim of transporting causal knowledge from s to s^* . More precisely, the available information is $P(y | do(a), X = x, S = s)$, $P(y, x, a | S = s)$ and $P(y, a, x | S = s^*)$, and the aim is to recover $P(y | do(a), S = s^*)$. The *transport formula*, which we prove in Annex E.1, addresses this task as follows.




Figure 5: Illustration of selection diagrams depicting differences between source and target populations: In (a), the two populations differ by their X distributions (so S points to X). In (b), the populations differ in how X depends on an unmeasured variable, represented by the open circle and the unmeasured variable distributions are the same. In (c), the populations differ in how X depends on A . In (d), the unmeasured confounder (bidirected arrow) between X and Y precludes transportability.

Theorem 5.2 (Transport formula) *The causal effect $P(y | do(a), S = s^*)$ can be transported from a population Π to a target population Π^* , if there exists a set X of observed pretreatment covariates that is **S -admissible**, i.e., when X blocks all paths from S to Y after deleting from the graph all arrows into A , that is, S satisfies $(S \perp\!\!\!\perp Y | X)_{G_{\bar{A}}}$. $G_{\bar{A}}$ represents the diagram after removing all arrows pointing at A . The transport formula is given by the weighting:*

$$P(y | do(a), S = s^*) = \sum_x P(y | do(a), X = x, S = s) P(x | S = s^*). \quad (13)$$

Note that S -admissibility can be seen as the graphical mirror of S -ignorability of Assumption 3.5. Then, the transport formula (13) can be seen as the counterpart of the identifiability formula (see (14)). Indeed, in our example, the trial data corresponds to $S = s$ and is unconfounded so that $P(y | do(a), x, S = s) = P(y | a, x, S = s)$.

In more complicated graphs, for instance involving post-treatment variables (selection biases) as illustrated in Figure 4 (c), domains may differ due to variables that are themselves causally affected by the treatment. The transport formula can then be different (weighting by $P^*(x | a)$ instead of $P^*(x)$). In complex setting, some automatic algorithmic solutions are available to determine if (non-parametric) identifiability holds (Correa et al., 2018; Tikka et al., 2019).

5.2 Complementarity of the PO and SCM frameworks

It appears that there is an imbalance in the focus of the two frameworks: identifiability is dealt with in two paragraphs in the PO framework and requires several pages in the SCM framework; as for the estimation, it takes a large part in the PO framework and is non-existent in the SCM framework.

In classical problems, both formulations lead to the same identifiability expressions, as we highlighted in this section: the backdoor criterion in Theorem 5.1 and the transport formula (13) correspond respectively to identifiability formulae (2) and (10) in the PO framework. Equation (2) is useful to make the connection between the PO and SCM frameworks, because with a discrete X , (2) becomes:

$$\mathbb{E}[\mu_{a,1}(X)] = \sum_x \mu_{a,1}(X)P(X=x). \quad (14)$$

Nevertheless, in more complex settings, the SCM framework with its elegant way of formalizing the problem, helps establishing identifiability. In addition, one important practical advantage of using SCMs is, subject to a good knowledge of the graph, to be able to select sets of variables that are sufficient to establish identifiability and, above all, to exclude variables that would bias the analysis. As we will see in Section 7 the DAGs offer a very convenient tool to discuss with clinicians and explicitly lay out conditional independences.

However, the SCM framework does not provide yet the user with a ready-to-use solution to estimate the causal effect and does not detail the properties of the associated estimators. For instance, to estimate the causal effect using the the transport formula (13), one can naturally use the g-estimator in (5), where the conditional mean of outcome on treatment and covariates are estimated based on the RCT data and then average over the covariate distribution from the observational data. But other estimators are available and not mentionned. In addition, the SCM framework related work does not cover the topics discussed in Section 4, which focuses on estimating the CATE and combining biased and unbiased estimators.

6 Software for combining RCT and observational data

6.1 Review of available implementations

A decisive point to bridge the gap between theory and practice is the availability of softwares. Practitioners use the methods with easily-available implementations, even when these methods have some shortcomings. In recent years, there have been more and more solutions for users interested in causal inference and causation, see Tikka and Karvanen (2017); Guo et al. (2018); Yao et al. (2020) for surveys. One can mention the toolboxes doWhy (Sharma et al., 2019), econML (Research, 2019), causalToolbox (Künzel et al., 2018). There are also many standalone packages. However, one can regret that many implementations are ‘one-shots’, i.e., associated with a single article and not pushed further.

Regarding the specific subject of this article, we present in Table 3 the implementations available about both identifiability and estimators. The implementations found are mostly research-dedicated tools made public rather than user-friendly packages. Note that no implementation of stratification (4) was found. Most implementations do not handle categorical variables, or handle only continuous outcome for example. In addition, the available implementations are often dedicated to

specific sampling designs and estimators are different from nested and non-nested framework. As a consequence, a new user has to pay attention to all of these practical – but fundamental – details.

Table 3: Inventory of publicly available code for generalization (top: software for identification; bottom: software for estimation).

Name	Method - Setting	Source & Reference
causaleffect	Identification and transportation of causal effects, e.g., conditional causal effect identification algorithm	R package on CRAN, Tikka and Karvanen (2017)
dosearch	Identification of causal effects from arbitrary observational and experimental probability distributions via do-calculus	R package on CRAN, Tikka et al. (2019)
Causal Fusion	Identifiability in data fusion framework, (Section 5)	Browser beta version upon request ?
ExtendingInferences	IPWS equation (3), g-formula equation (S5) - Nested AIPSW (S7) - Nested Continuous outcome	R code on GitHub, Dahabreh et al. (2020)
generalize	IPSW equation (3), TMLE (Section 3.2.4)	R package on GitHub Ackerman et al. (2020)
genRCT	IPSW equation (S3) - Nested, calibration weighting (Section 3.2.4) Continuous and binary outcome	R package available upon request Dong et al. (2020)
IntegrativeHTE	Integrative HTE (Section 4.2)	R package on GitHub, Yang et al. (2020)
IntegrativeHTEcf	Includes confounding functions (Section 4.2)	R package on GitHub, Yang et al. (2020)
generalizing	SCM with probabilistic graphical model for Bayesian inference Binary outcome	R package on GitHub, Cinelli and Pearl (2020)
RemovingHiddenConfounding	Unmeasured confounder (Section 4.2)	R package on GitHub, Kallus et al. (2018)

6.2 Example of usage

In this part we detail an identifiability question. Implementation examples for the nested case are presented in examples D.3.1 and D.3.2.

Identifiability queries The R packages `causaleffect` (Tikka and Karvanen, 2017) and `dosearch` (Tikka et al., 2019) can be used for causal effect identification, with the later handling transportability, selection bias and missing values (bivariate) issues simultaneously. In this package, the `dosearch` function takes the observable distributions, a query, and a semi-Markovian causal graph as the input and outputs a formula for the query over the input distributions, or decides that it is not identifiable. It is based on a search algorithm that directly applies the rules of do-calculus. Their general identification procedure is not necessarily complete given an arbitrary query and an arbitrary set of input distributions. In order to retrieve the backdoor criterion in theorem 5.1, one can write:

```

1 data <- "P(Y, X, Z)"
2 query <- "P(Y|do(X))"
3 graph <- "X -> Y
4           Z -> X
5           Z -> Y"
6 dosearch(data, query, graph)

1 $identifiable
2 [1] TRUE
3 $formula
4 [1] "[sum_{Z} [p(Z)*p(Y|X,Z)]]"

```

6.3 Simulation study of the main approaches

This part presents simulations results to illustrate the different estimators introduced and their behavior under several mis-specifications patterns. The code to reproduce the results is available on Gitlab³. Note that except for the calibration weighting, all the estimators are implemented by the authors to correspond exactly to the formulae introduced in the review (IPSW and IPSW.norm (3), stratification (4), g-formula (5), and AIPSW (8)).

Well-specified models We consider the framework of Section 3 where the observational study contains neither outcome nor treatment and the aim is to estimate the causal effect on the target population. We use similar simulations as in Dong et al. (2020), where four covariates are generated independently as with $X_j \sim \mathcal{N}(1, 1)$ for each $j = 1, \dots, 4$. The trial selection scores are defined using a logistic regression model:

$$\text{logit } \{\pi_S(X)\} = -2.5 - 0.5 X_1 - 0.3 X_2 - 0.5 X_3 - 0.4 X_4. \quad (15)$$

The outcome is generated according to a linear model:


$$Y(a) = -100 + 27.4 a X_1 + 13.7 X_2 + 13.7 X_3 + 13.7 X_4 + \epsilon \text{ with } \epsilon \sim \mathcal{N}(0, 1). \quad (16)$$

This outcome model implies a target population ATE of $\tau = 27.4 \mathbb{E}[X_1] = 27.4$. Note that the sample selection ($S = 1$) in (15) is biased toward lower values of X_1 and consequently toward lower treatment effect.

To generate a non nested trial, we proceed as followed. First a sample of size 50,000 is drawn from the covariate distribution. From this sample, the selection model (15) is applied which leads

³<https://gitlab.inria.fr/misscausal/combine-rct-rwd-review>

Figure 6: Well-specified model
 Estimated ATE with the inverse propensity of sampling weighting with and without weights normalization (IPSW and IPSW.norm; Section 3.2.1), stratification (with 10 strata; Section 3.2.1), g-formula (Section 3.2.2), calibration weighting (CW; Section. 3.2.3), and augmented IPSW (AIPSW; Section 3.2.4) over 100 simulations.



to an RCT sample of size $n \sim 1000$. Then, the treatment is generated according to a Bernoulli distribution with probability equals to 0.5, $e_1(x) = e_1 = 0.5$. Finally the outcome is generated according to (16). The observational sample is obtained by drawing a new sample of size $m = 10,000$ from the distribution of the covariates.

Figure 6 presents estimated ATE with the inverse propensity of sampling weighting with and without weights normalization (IPSW and IPSW.norm; Section 3.2.1), stratification (with 10 strata; Section 3.2.1), g-formula (Section 3.2.2), calibration weighting (CW; Section. 3.2.3), and augmented IPSW (AIPSW; Section 3.2.4) over 100 simulations. All estimators are implemented by the authors, except for the CW implementation which comes from Dong et al. (2020). The true ATE is represented with a dash line. The ATE estimated only with the RCT sample is also displayed as a baseline. As expected it is biased downward (its mean is equal to 14.24) as the distribution of the covariates and in particular the treatment effect modifiers such as X_1 is not the same in the trial sample and in the population (as illustrated in Table 17 in Appendix F). Note that all the estimators are unbiased which is expected. The variability of the two IPSW estimators are larger than the others. The number of strata in the stratification estimator plays an important role. As shown in Figure 18, the results are biased when the number of strata is smaller than 10.

Mis-specification of sampling propensity score or outcome model To study the impact of mis-specification of the sampling propensity score model, we generate the RCT selection according to the model

$$\text{logit} \{ \pi_S(X) \} = -2.5 - 0.5 e^{X_1} - 0.3 e^{X_2} - 0.5 e^{X_3} - 0.4 e^{X_4} + 3,$$

and outcome according to the model

$$Y(a) = -100 + 27.4 a X_1 X_2 + 13.7 X_2 + 13.7 X_3 + 13.7 X_4 + \epsilon.$$





Figure 7: Mis-specified models Estimated ATE when RCT and/or outcome models are mis-specified. Estimators used being IPSW (IPSW and IPSW.norm; Section 3.2.1), stratification (with 10 strata; Section 3.2.1), g-formula (Section 3.2.2), calibration weighting (CW; Section 3.2.3), and augmented IPSW (AIPSW; Section 3.2.4) over 100 simulations.


The analysis is then performed using classical logistic and linear estimators. As shown in Figure 7, when the sampling propensity score model is mis-specified, the IPSW estimators are biased; whereas when the outcome model is mis-specified, the g-estimator is biased. In both settings, the double robust estimator (AIPSW) is unbiased and robust to mis-specification. In the case where both models are mis-specified, all estimators are biased except the CW estimator. This demonstrates some robust properties of calibration against slight model mis-specification.

Stronger distributional shift The above sampling propensity score model implies a weak covariate shift between the RCT sample and the observational sample. A stronger shift can be obtained, at least on covariate X_1 , swapping the coefficient $-0.5X_1$ with $-1.5X_1$. Figure 8 shows that the variance of the weighted and CW estimators have increased in the setting with a stronger covariate shift.

Impact of the treatment-effect modifiers In this part, we consider a heterogeneous treatment effect setting where X_1 impacts the RCT sampling while also being a treatment effect modifier. We consider the IPSW estimator and its variations without using X_1 (labeled as IPSW.without.X1) and using only X_1 (labeled as IPSW.X1). As shown in Figure 9, IPSW.X1 is still unbiased when using only X_1 in the sampling propensity score estimation, as it is the only covariate being the



(a) Distributional shift on X_1



(a) Simulation results

Figure 8: Weak versus strong distributional shift between experimental and observational data (a) illustrates the different strength of the distributional shifts, and (c) presents estimated ATE when RCT is weakly or strongly shifted from the target population distribution. Estimators used being IPSW (IPSW and IPSW.norm; Section 3.2.1), stratification (with 10 strata; Section 3.2.1), g-formula (Section 3.2.2), calibration weighting (CW; Section. 3.2.3), and augmented IPSW (AIPSW; Section 3.2.4) over 100 simulations.


treatment effect modifier. However, if X_1 is missing, the resulting estimator IPSW.without.X1 is strongly biased. Therefore, it is important to include all variables that affect both sampling and outcome to adjust for bias. We also conjecture that including outcome predictors that do not affect sampling may increase the efficiency of the estimator.

Note also that if the treatment effect was homogeneous (does not depend on X_1), then the estimated ATE on the RCT would be unbiased (as shown Figure 19 in Appendix F.3) so in this setting there is no need to use the observational data and associated methods to transport the ATE from the trial to the target population.

7 Application: Effect of Tranexamic Acid

To illustrate the methodological question of combining experimental and observational data and demonstrate how to apply some of the previously discussed methods, we consider a currently open medical question about major trauma patients. Major trauma denotes injuries that endanger the life or the functional integrity of a person. The World Health Organization (WHO) has recently shown that major trauma including road-traffic accidents, interpersonal violence, falls, etc. remains a world-wide public-health challenge and major source of mortality and handicap (Leigh et al., 2018). We focus on trauma patients suffering from a traumatic brain injury (TBI). TBI is a sudden

Figure 9: Impact of the treatment-effect modifiers
 Estimated ATE when IPSW estimator includes all covariates, only X_1 , or all covariates except X_1 (IPSW; Section 3.2.1), with g-formula (Section 3.2.2) presented as a control, over 100 simulations. Simulations are still performed with (15) for RCT eligibility and (16) for outcome modeling.



damage to the brain caused by a blow or jolt to the head and can lead to intracranial bleeding that can be observed on a computed tomography (CT) scan. Ongoing intracranial bleeding can lead to raised intracranial pressure, brain herniation, and death. Over 10 million people are killed or hospitalized worldwide because of TBI each year (Dewan et al., 2012). Tranexamic acid (TXA) is an antifibrinolytic agent that limits excessive bleeding, commonly given to surgical patients. Previous clinical trials showed that TXA decreases mortality in patients with traumatic *extracranial* bleeding (Shakur-Still et al., 2009). Such prior result raises the possibility that it might also be effective in TBI, because *intracranial* hemorrhage is common in TBI patients. Therefore the question here is to assess the potential decrease of mortality in patients with intracranial bleeding when using TXA.

To answer this question, we have at disposal both a RCT, named CRASH-3, and an observational study, the Traumabase. Both data have previously been analyzed separately in CRASH-3 (2019); ? (for the RCT) and in Mayer et al. (2020) (for the observational study) and the medical teams of both studies want to share their respective data to answer both medical and methodological questions. Such initiatives allow to confront and combine the evidence obtained from the observational data set using causal inference estimators with the evidence obtained from the CRASH-3 randomized controlled trial. In particular it allows to:

- Benchmark the observational study with the RCT. As the treatment effect was estimated with a recent doubly robust estimator handling missing data, its adoption by the community could be reinforced using the RCT's result to validate the estimated effect.
- Assess the generalizability and transportability methods, considering the Traumabase as the target population, and confront the estimators presented in this review to a real situation.

We will first present the two data sources, treatment effect analyses and findings from these, before turning to the combined analysis in Section 7.3. The code to reproduce all analyses for this

application is available on GitLab⁴, even if the medical data cannot be publicly shared for privacy reasons.

7.1 The observational data: Traumabase

To improve decisions and patient care in emergency departments, the Traumabase group, which comprises 23 French Trauma centers, collects detailed clinical data from the scene of the accident to the release from the hospital. The resulting database, called the Traumabase, comprises 23,000 trauma admissions to date, and is continually updated. The data are of unique granularity and size in Europe. However, they are highly heterogeneous, with both categorical – sex, type of illness, ... – and quantitative – blood pressure, hemoglobin level, ... – features, multiple sources, and many missing data (in fact 98% of the individuals have missing values). The cause of missing information is also coded, such as technical hurdles with the measurement, or impossibility due to the severity of the patient’s state. The Traumabase currently comprises around 8,270 patients suffering from TBI. A first study was performed to assess the effect of TXA on mortality for traumatic brain injury patients from this observational registry (Mayer et al., 2020). More precisely, the treatment variable is the administration of TXA during pre-hospital care or on admission to a Trauma Center⁵ and considered to have occurred within three hours of the initial trauma. In the current Traumabase, TXA is administered to roughly 8.2% of patients suffering from TBI, and 20% die before the end of their hospital stay. Notably, mortality is much higher among patients who receive TXA than those who do not (30% vs. 14%). This situation is a classical example of confounding bias: the effect arises because patients who appear to be in more severe state are more likely to be administered TXA and are also more likely to die, with or without the treatment.

Before turning to the causal analysis on these data, we first discuss an important practical aspect, namely missing values, and how we handle them in the subsequent analyses.

7.1.1 Missing values

The problem of missing values is ubiquitous in data-analysis practice and particularly present with observational data, as they are not necessarily collected for research purposes. The Traumabase is a high-quality dataset but, nevertheless, missing values occur. Figure 10 represents the percentage of missing values for the covariates selected by the medical doctors from the Traumabase. It varies from 0 to nearly 60% for some features. In addition, there are different codes for missing values giving hints on the reason of their occurrence, resp not available (NA), impossible (imp), not made (NM), etc. Some of these values can be seen as missing completely at random (MCAR), the information has not been recorded simply because the form was not filled out, but they can be informative and missing not at random (MNAR), for instance when the state of the patient is such that it was impossible to take a measurement.

⁴<https://gitlab.inria.fr/misscausal/combine-rct-rwd-review>

⁵More precisely, to the resuscitation room of a hospital equipped to treat major trauma patients.




Figure 10: **Missing values** Percentage of missing values for a subset of Traumabase variables relevant for traumatic brain injury. Different encodings of missing values are available such as: *NA* (not available), but also *not informed*, *not made*, *not applicable*, *impossible*.

There is an abundant literature available to deal with missing values in a general context and Mayer et al. (2019) identify more than 150 R (R Core Team, 2018) packages available on the topic. Missing values add a layer of complexity to conducting causal analyses as they require coupling conventional hypotheses of causal-effect identifiability in the complete case with hypotheses about the mechanism that generated the missing data (Rubin, 1976), or defining new hypotheses, to establish conditions of causal effect identifiability with missing data. Mayer et al. (2020) survey available works, classify the methods in three families that differ with respect to the different assumptions and provide associated estimators to estimate the ATE from an observational data set with missing values in the covariates. More precisely, they advocate the use of multiple imputation (van Buuren, 2018) by IPW or doubly robust estimators when missing values can be considered to be missing (completely) at random (M(C)AR) and the classical unconfoundedness assumption (Assump. 4.1) holds (Seaman and White, 2014). As an alternative, they recommend using a doubly robust estimator adapted to missing values, i.e., that makes use of random forests with a missing incorporate in attributes splitting criterion (Twala et al., 2008; Josse et al., 2019) to estimate the generalized propensity scores (Rosenbaum and Rubin, 1984) and the regression function with missing values⁶; this approach does not require a particular missing values mechanism but an adapted unconfoundedness hypothesis with missing data. Finally, when covariates can be seen as noisy incomplete proxies of true confounders, latent variable models can be a solution to estimate causal effect with missing values (Kallus et al., 2018; Louizos et al., 2017).

⁶This doubly robust method is implemented in the R package grf (Athey et al., 2019).

7.1.2 Covariate adjustment

Since the Traumabase is an observational registry, straightforward treatment effect estimation on these data is not possible due to confounding. The causal graph in Figure 11 is the result of a two-stage Delphi method (Linstone and Turoff, 1975) in which six anesthetists and resuscitators specialized in critical care—and therefore familiar with the allocation process for TXA—first selected covariates related to either treatment or outcome or both, and second classified these covariates into confounders and predictors of only treatment or outcome. Even though it is not possible to test for unobserved confounding, this Delphi procedure is an attempt to gather as much expert knowledge about the studied question as possible to manually identify possible confounders and qualitatively assess the plausibility of the unconfoundedness assumption. Note that this approach is an explicit example where we leverage the advantages of the SCM and PO frameworks: the causal graph helps to select relevant variables during the conception phase of the study, and the treatment effect analysis uses different estimation methods from the PO framework.




Figure 11: **Causal graph** representing treatment, outcome, confounders and other predictors of outcome (Figure generated using DAGitty (Textor et al., 2011); **NAs** indicates variables that have missing values).

7.1.3 Results

We adjust for confounding using the identified confounders and use additional predictors for the outcome model (Figure 11), to estimate the direct causal effect of TXA on 28-day intra-hospital TBI-related mortality and on all cause intra-hospital mortality among traumatic brain injury patients

(for the latter see Mayer et al. (2020)). The two methods presented here for handling missing values are multiple imputation via chained equations (MICE) (van Buuren, 2018) and missing incorporated in attributes (MIA) (Twala et al., 2008); the treatment effect estimation is then performed using either logistic regressions or generalized random forests (Athey et al., 2019), denoted respectively by *GLM* and *GRF* in Table 4.⁷ The doubly robust results (AIPW) in Table 4 show that from this study there is no evidence for an effect of TXA on mortality of TBI patients. However, when considering the IPW estimations, the conclusion differs in that they indicate a deleterious effect of the drug for almost all subgroups considered, for both definitions of the outcome. These findings might be due to inaccurate estimations of the propensity scores used for the reweighting, for instance due to possible model mis-specification for the parametric approach and to insufficient sample sizes or machine learning regulation bias for the non-parametric random forest approach. In the non-parametric case, the doubly robust approach can compensate the slow convergence of the non-parametric propensity model estimations with the outcome model estimations and correct for the regularization bias that comes from fitting predictive models using machine learning, while there is no such bias correction for the IPW approach and in case of insufficient the latter would also require additional samples to estimate the propensity scores sufficiently well. Additionally, note the large variability of the parametric IPW and AIPW estimators (*GLM*) which could also support the remark on possible model mis-specification.⁸ In such a situation, the possibility to transport the treatment effect from the RCT is also a step to comfort the results.

The AIPW findings on the data—obtained prior to the publication of CRASH-3—are consistent with the main conclusion of the CRASH-3 study (no effect). Results are presented in Table 4). As patients can be stratified on the trauma severity, analysis on sub-strata can also be performed, but are only reported in the Appendix G.

Table 4: ATE estimations from the Traumabase for TBI-related 28-day mortality. Red cells conclude on deteriorating effect, white cells conclude on no effect. Additional results can be found in Table 7 in the Appendix G

	Multiple imputation (MICE)				MIA		Unadjusted ATE $\times 10^2$
	IPW (95% CI) $\times 10^2$		AIPW (95% CI) $\times 10^2$		IPW (95% CI) $\times 10^2$	AIPW (95% CI) $\times 10^2$	
	GLM	GRF	GLM	GRF			
Total (n = 8248)	15 (6.8, 23)	11 (6.0, 16)	3.4 (-9.0, 16)	-0.1 (-4.7, 4.4)	9.3 (4.0, 15)	-0.4 (-5.2, 4.4)	16

7.2 The RCT: CRASH-3

A total of 175 hospitals in 29 different countries participated to the randomized and placebo-controlled trial, called CRASH-3 (Dewan et al., 2012), where adults with TBI suffering only from intracranial bleeding, i.e., without major extracranial bleeding, were randomly administrated TXA (CRASH-3, 2019; ?). The inclusion criteria of the trial are patients with a Glasgow Coma Scale

⁷Note that another method for handling the missing values could theoretically be used, namely EM for logistic regression (Jiang et al., 2020). In this application however, this method is not (yet) adapted due to the large number of mixed covariates.

⁸Larger variance of the IPW is often observed when comparing to AIPW, independently of the nuisance parameter approach.

(GCS)⁹ score of 12 or lower or any intracranial bleeding on CT scan, and no major extracranial bleeding, leading to 9202 patients (which is uncommonly large for a medical RCT). We provide a DAG summarizing the trial selection and addition regressors present in CRASH-3 of the outcome in Figure 12.




Figure 12: **Causal graph** representing treatment, outcome, inclusion criteria with S and other predictors of outcome (Figure generated using the Causal Fusion software presented in Section 6 from ?).

The primary outcome studied is head-injury-related death (not simply death) in hospital within 28 days of injury in patients treated within 3 hours of injury. Six covariates are present at baseline, with age, sex, time since injury, systolic blood pressure, Glasgow Coma Scale score (GCS), and pupil reaction. The study concludes that the risk of head-injury-related death was 18.5% in the TXA group versus 19.8% in the placebo group (855 vs 892 events; risk ratio [RR] 0.94 [95% CI 0.86 - 1.02]), i.e, there is no effect of TXA on mortality when considering the total sample. The rest of the data analysis focuses on the effect on the total sample as well.

7.3 Transporting the ATE on the observational data

With the two separate analyses in mind, we are now ready to tackle the combined analysis, namely the generalization from the RCT results to the target population defined by the observational Traumabase data.

7.3.1 Common covariates description

In the following, we discuss common variables definition, outcome, treatment, and designs in order to leverage both sources of information. We recall the causal question of interest: “What is the effect of the TXA on brain-injury death on patients suffering from TBI?” This part is important for the harmonization of the study protocol.

Treatment exposure. The treatment protocol of CRASH-3 frames the timing and mean of administration precisely (a first dose given by intravenous injection shortly after randomization,

⁹The Glasgow Coma Scale (GCS) is a neurological scale which aims to assess a person’s consciousness. The lower the score, the higher the gravity of the trauma.

i.e., within 8 (resp. 3) hours of the accident, and a maintenance dose given afterwards (Dewan et al., 2012)). For consistency with the original CRASH-3 study described above, we also only keep observations from the RCT with administration within 3 hours. The Traumabase study being a retrospective analysis, this level of granularity concerning TXA is not available. Neither the exact timing, nor the type of administration are specified for patients who received the drug. However, the expert committee agreed that the assumption of treatment within 3 hours of the accident is very likely since this drug is administered in pre-hospital phase or within the first 30 minutes at the hospital.

Outcome of interest. The CRASH-3 trial defined its primary outcome as head injury-related death in hospital within 28 days of injury. For the Traumabase data we also look at death in hospital within 28 days but with a wider range of possible causes of death, namely TBI, brain death, multiple organ failure, brain death, or withdrawal of life-sustaining therapy. This slightly different definition for the Traumabase outcome allows to obtain medically similar outcomes despite different data collection processes.

Multi-centered design. Both studies are multi-centered, but while the Traumabase is a French registry with 20 participating Trauma Centers, the CRASH-3 trial enrolled patients in various countries on different continents. While this large spectrum of participating centers is likely to contribute to external validity of the CRASH-3 trial, it should be noted that more than 65% of the patients included were from developing countries; regions of the world that differ from developed countries by a prolonged pre-hospital care period, limited access to brain imaging tests and neurosurgery within short periods of time, and the absence of expert centers for heavy trauma and neuro-intensive care. Thus, on top of the restrictive inclusion criteria of the RCT, this aspect of large heterogeneity in the participating Trauma centers motivates the combination both studies to estimate the effect for a population with access to a specific high level of care, here represented by the French Trauma centers.

Covariates accounting for trial eligibility. In total, four criteria depending on five variables determined inclusion in the CRASH-3 trial: age (only adults were eligible), presence of TBI (defined as presence of intracranial bleeding on the CT scan, or a GCS of less than 13 in the case of no available CT scan), absence of major extracranial bleeding (defined explicitly in CRASH-3 and defined via the number of packed red blood cells transfused in the first 6 hours of admission or by colloid injection in the Traumabase), and delay of less than 8 hours (later reduced to 3 hours) between the injury and the randomization. The necessary variables are also available in the Traumabase, either exactly or in form of proxies, which allows the estimation of the trial inclusion model on the combined data.

Additional covariates. Note that other covariates are available in both data sets, while not responsible of trial inclusion according to CRASH-3 investigators. But as this could still be covariates moderating the outcome and treatment effect, we included them. According to the two data sets, we could add three of them: sex (binary), systolic blood pressure (continuous), and pupils reactivity (categorical, ranging from 0 to 2, being the number of active pupils). Note that these three covariates are all mentioned in the baseline of CRASH3 results (CRASH-3, 2019), arguing that they should impact the outcome.

7.3.2 Descriptive analyses

Missing values. First, note that the RCT contains almost no missing values, whereas the variables for determining eligibility in the observational data contains important fractions of missing values, as shown in Table 5, while the sample sizes of the two data are similar, see Table 6. This requires a modification of the methods introduced in this review and illustrated in the simulations section in order to account for these missing values. For correctly estimating the trial inclusion model and the outcome models (Section 3.2), we need to handle the missing values in the covariates, especially in the observational data¹⁰. This modification consists in two alternative estimation strategies for fitting the trial inclusion and outcome models:

- Logistic regression with incomplete covariates using an expectation maximization (Dempster et al., 1977). A computationally efficient variant of this method using stochastic approximation is implemented in the R package `misaem` (Jiang et al., 2020).
- Generalized regression forest with missing incorporated in attributes (Twala et al., 2008; Josse et al., 2019). This method is implemented in the R package `grf` (Tibshirani et al., 2020).

Table 5: Percentage of missing values in each covariates for the Traumabase and CRASH-3.

	MajorExtracranial	Age	Glasgow initial	SystolicBloodPressure	Sex	Pupil Reactivity
CRASH-3	0	0	0.69	0.25	<0.1	<0.1
Traumabase	0	0.27	2.0	29	0.76	2.0

Table 6: Sample sizes for both studies.

	Traumabase			CRASH-3		
	m	#treated	#death	n	#treated	#death
Total (within 3 hours)	8248	683	1411	9168	4632	1745

Distribution shift. There are different ways of assessing the shift between the distributions of the two studies, for instance by univariate comparisons. We provide a simplified comparison of the means of the covariates between the treatment groups of the two studies in Figure 13. This graph illustrates again the fundamental difference between the two studies, namely the treatment bias in the observational study and the balanced treatment groups in the RCT. Another representation of the distribution shift is presented in the Appendix G with histograms (Figures 20–24).

¹⁰If we assumed the missing values being missing completely at random (MCAR), we could “throw away” the incomplete observations and perform the analyses on the complete observations, but this would reduce the total sample size by 917. And as explained in Section 7.1, the MCAR assumption is not plausible for the present observational data, thus such a *complete case analysis* would be biased.

	majorExtracranial	Glasgow.initial	age	pupilReact_num	systolicBloodPressure	sexe	TBI_Death
Control.Observational	0.65	10.81	43.29	1.67	130.18	0.22	0.16
Treated.Observational	0.99	8.42	41.73	1.27	100.14	0.33	0.32
Control.RCT	0	9.58	41.9	1.65	129.64	0.2	0.2
Treated.RCT	0	9.62	41.75	1.64	130.41	0.19	0.18

Figure 13: **Distributional shift** and difference in terms of univariate means of the trial inclusion criteria (red: group mean greater than overall mean, blue: group mean less than overall mean, white: no significant difference with overall mean, numeric values: group mean (resp. proportion for binary variables). Graph obtained with the `cateds` function of the `FactoMineR` package (Lê et al., 2008).

7.3.3 Analyses

Notations and estimator details. We use two consistent ATE estimators from the solely CRASH-3 data:

- `Difference_in_mean`: the difference in mean estimator (Section A);
- `Difference_in_condmean_ols` the difference in conditional mean fitting an outcome model with an OLS.

To transport the ATE to the target population, we apply the estimators discussed in this review (with the additional handling of the missing values as outlined in the previous section), namely:

- IPSW (3): with sampling propensities estimated via logistic regression (`EM_IPSW_glm`) or via generalized random forest (`MIA_IPSW_grf`);
- normalized IPSW (3): with sampling propensities estimated via logistic regression (`EM_IPSW.norm_glm`) or via generalized random forest (`MIA_IPSW.norm_grf`);
- G-formula (5): with outcome models estimated via logistic regression (`EM_G-formula_glm`) or via generalized random forest (`MIA_G-formula_grf`);
- AIPSW (8): with sampling propensities and outcome models estimated via logistic regression (`EM_AIPSW_glm`) or via generalized random forest (`MIA_AIPSW_grf`).

The confidence intervals of those estimators are computed with a stratified bootstrap in the RCT and the observational data set in order to maintain the ratio of relative size of the two studies (with 100 bootstrap samples). Note that the Calibration Weighting (CW) estimators are not used in this

analysis as it would require a specific adaptation in the case of the missing values. Propensity score methods and outcome modelling methods are more straightforward to adapt to the missing data situation with off-the-shelf computational tools. Since this topic combines missing data assumptions with the specific assumptions for the goal of generalization, the specific validity domain for each estimator remains an open-research work.

We also present the estimators for the observational study applied solely on the Traumabase data. For details about the derivation and properties of these estimators applied on incomplete observations we refer to Mayer et al. (2020):

- AIPW with nuisance parameters via logistic regression (**AIPW_glm**),
- AIPW with nuisance parameters via generalized random forest (**AIPW_grf**).

Since AIPW combined with either missing incorporated in attributes (MIA) or multiple imputation (MICE) is recommended when analyzing observational data, these are the estimators kept in this analysis. These estimators are built upon the unconfoundedness assumption and as outlined in Section 7.1, the list of identified confounders comprises 17 variables, complemented with 21 variables predictive of the outcome but not related to the treatment. Hence the results of the observational study depend on these 38 variables while the generalization results are using a different set of variables, namely three of the five variables that determine treatment eligibility¹¹ and the additional three baseline covariates susceptible to moderate outcome and treatment effect.

Final results. As we can observe on Figure 14, the generalization from the RCT to the target population using all the observations from both data sets, presents certain discrepancies with respect to the two previous studies. On the one hand, half the generalization estimators support the CRASH-3 conclusion about the treatment effect: no significant effect. On the other hand, some estimators support a deleterious treatment effect (corresponding to a positive ATE). Note that the AIPW ATE estimations from the solely Traumabase data do not reject the null hypothesis of no treatment effect. The analysis can also be performed on a imputed Traumabase data set, the corresponding results are presented in appendix on Figure 29. Note that these results are to be interpreted carefully due to the potential impact of missingness on the performance of the chosen estimators. For instance, note the large confidence intervals for the GRF estimators that are likely to be due to the imbalanced proportions of missing values in the RCT and the observational data.

Here we present the results transported onto the total TBI Traumabase population, but the CRASH-3 study focuses on subgroups of patients (mild and moderate patients) for which a positive effect of the tranexamic acid is measured. The generalization of the CRASH-3 findings onto this subgroup in the Traumabase raises multiple methodological issues that still need to be addressed in future works and that we detail in the Appendix G.

Overall this data analysis highlights practical limitations that can be encountered when combining two different data sets: the need for a good understanding of the common covariates, exposure, and outcome of interest before combining the data sets, different missing data patterns, and poor overlap when considering specific target (sub-)populations.

¹¹We filter the CRASH-3 population using the remaining two variables, TBI and delay between injury and randomization, therefore these variables are removed for the remainder of the analysis.




Figure 14: **Juxtaposition of different estimation results** with ATE estimators computed on the Traumabase (observational data set), on the CRASH-3 trial (RCT), and transported from CRASH-3 to the Traumabase target population. All the observations are used. Number of variables used in each context is given in the legend.

8 Summary, recommendations, and shortcomings

Combining observational data and RCTs can improve many aspects of causal inference, from increased statistical power to better external validity. Questions on external validity arise as soon as there are heterogeneities in the populations under study, whether these be heterogeneities of a treatment effect or simply heterogeneities of outcome. And yet, most effects of interest have some form of heterogeneity: mortality rates increase with age, different populations are exposed to different risks, have different comorbidities, etc. Generalization challenges may arise even for homogeneous treatment effects, as developed by Cinelli and Pearl (2020) in a thought experiment on the Russian Roulette: its treatment is homogeneous –for every game, there is one chance out of six of dying– but the average treatment effect may differ in across populations is the chance of dying of other, independent, reasons differ. To come to clear conclusions, an RCT is often run on a more homogeneous subpopulation. The definition of the corresponding inclusion criteria is then crucial. Observational data can help defining these inclusion criteria, and the methods reviewed in the present paper, combining the RCT and observational data, can extrapolate the results of the RCT to different populations, supplementing other RCTs with different inclusion criteria. These statistical methods are crucial to explore potential heterogeneities left aside in the RCT. A plethora of settings is associated with a variety of identifiability criteria and estimators. Yet, whether in

the potential-outcomes or in the structural causal models framework, the number of methods for combining experimental and observational data with treatment and outcome is still limited.

Identifiability: which data to answer our question? Domain expertise can be used to postulate a causal graph: a directed acyclic graph representing the mechanisms (as Figure 12). The SCM framework is then convenient to see whether the question at hand can be formulated in an identifiable way. It offers a principled way of selecting variables needed for identification of the causal effect and to avoid biased causal effects, such as conditioning on the wrong covariates.

Without such an approach, identifiability claims are limited and the recommendation is often to include as many as possible to be sure to avoid violation of any assumption: “it is probably best to include as many outcome predictors as possible in regression models for the expectation of the outcome or the probability of trial participation” (?).

Selecting a small number of covariate can help reducing the variance of estimators. Some developments in the broad settings of causal inference –and not dedicated to fusion of experimental and observational data– use causal graphs to select between different adjustment sets to get estimators with smaller variance (Rotnitzky and Smucler, 2019; Witte et al., 2020; Guo and Perković, 2020).

Challenges in the SCM framework. A challenge in the use of the SCM framework is that, to fully exploit its potential, it requires a graph, much more detailed than the one we could establish for the Traumabase analysis. The consortium of clinicians managed to formulate groups of variables to identify confounders but not the precise links between all (or a set of) these variables.

Also, the current developments of the SCM framework do not mention treatment effect heterogeneity neither explicitly the differences between nested and non-nested designs. Finally, the corresponding literature is typically useful to the practitioner for variable selection, but lacks estimators that can be readily instantiated on the data –with the notable exception of Cinelli and Pearl (2020) in a Bayesian setting.

Estimation: what are the treatment effects in wider populations? For causal effects on a combination of experimental and observational data, available estimators mostly lie in the potential outcome framework. Estimation typically relies on a propensity score –modeling trial participation– or outcome-regression models. Weakly-parametric models such as machine-learning estimators or estimators coping with missing values can be used both, enabling a large class of functional forms. A different approach is that of calibration weighting, which rather relies on balancing the functions of the covariates. Theoretical properties of different estimators, and their behavior in our experimental study, outline practical recommendations. Calibration weighting and doubly-robust methods (AIPSW) are more robust to misspecification. Calibration weighting is more robust than doubly-robust methods which break down if both the propensity-score and the outcome regression model are mis-specified (Figure 7). But, on the other hand, calibration weighting suffers more variance in case of large shifts of the causal effect between the RCT and the observational data (Figure 8) and is expected to break down with large dimension of covariates when there is no natural function to match. Modeling the probability of trial participation via stratification can help if this probability takes very high or very low value; however definition of the strata –in particular their number– can be important. Too little biases towards the RCT (Figure 18), akin to a typical bias-variance compromise. Any model of trial participation must capture correctly the treatment-effect modifier, in the sense of heterogeneous treatment effect (Figure 9). Conversely it can be to capture its effect, ignoring other covariates. As a consequence, the data collection and

modeling should focus of on covariate likely to modulate the treatment effect. For reliable causal conclusions, a good knowledge of the data and the mechanisms at hand is crucial. Indeed, not only is it important to formulate good models of the response or the selection biases, but also different settings, i.e, the nested and non nested design, lead to different estimators (this is often implicit in the publications). A few estimators are readily-available, but they must be used for the particular design they were conceived for.

Challenges to handle missing values. We saw the need to account for missing values. Missing values occur more often in observational data since RCT typically deploy significant efforts to avoid them, but they are not immune to participants missing scheduled visits or completely dropping out from the study. The literature for RCT mainly focuses on missing outcome data and calls for sensitivity analysis giving that available strategies (weighting, multiple imputation) rely on untestable assumptions on the missing values mechanism (Carpenter and Kenward, 2007; Little et al., 2012; Kenward, 2013; O’Kelly and Ratitch, 2014; Li and Stuart, 2019; Cro et al., 2020). In observational studies, many methods are available relying on different assumptions either on the missing values mechanism or on the identifiability conditions of the causal effect (Mayer et al., 2020).

Missing values may lead to subtle biases in the inferences, as they are seldom uniformly distributed across both datasets—missing more in one than in the other. Further research is needed on identifiability conditions and estimators in these setting, to better understand the scope of each estimator. Given its good properties in the complete data case, it could be interesting to study the calibration estimator with missing values.

An extreme case of missing information arises from the fact that observational data and clinical trials are seldom collected to be analyzed jointly. As such, they typically measure different covariates. The common practice to analyze these data jointly consists in only considering the covariates present in both data. However, throwing away covariates leads to lost opportunities to characterize better confounding effects and variables responsible for trial eligibility. Nguyen et al. (2017) and Nguyen et al. (2018) use extra covariates present in the experimental data to assess the sensitivity of the estimated treatment effect to an unobserved treatment effect moderator that interacts with treatment in influencing the outcome when generalizing from an RCT to a target population.

Discrepancies between RCTs and observational data. The analysis of real-world data allows to apply the methods and to assess their feasibility. This analysis is pointing towards moderate external validity of the RCT since the results where the ATE is generalized are not entirely concordant with the RCT. Notably, the study using only the observational data supports the results from the RCT. Which analysis to trust depends on the assumptions we are willing to make, either the transportability assumptions or the unconfoundedness assumptions and whether the variables have been selected appropriately. However, caution is needed when interpreting the presented results, since the impact of the missing values when combining RCT and observational data requires further research as mentioned above. In addition, the fact that outcome, treatment and covariates are comparable is of major importance (Lodi et al., 2019) and remains a challenging question when it comes to data fusion. The observed discrepancies support a possible deleterious effect of the treatment, and as soon as the treatment effect depends on the timing, another explanation of the difference could come from slightly different time-to-treatment in both data sets. Such considerations, beyond others, have to be investigated jointly with the clinicians. Finally, the results are presented for the total population but in this application, the clinicians are more interested in assessing treatment

effects on specific strata (see Appendix G for more details). However, there are issues to be solved before answering their request. Indeed, when considering certain strata in our application we are facing the issue of violated positivity, which leads to a non-transportable treatment effect on the strata of interest: mild and moderate patients. Therefore, further discussions and analyses with the medical expert committee are necessary to re-define a target population of interest on which generalization is possible and medically relevant. As it is generally the case, beyond methodological and theoretical guarantees, a major step before applying a set of methods is to clearly state the causal question and estimand(s) and the associated identifiability requirements.

Acknowledgment

This work is initiated by a SAMSI working group jointly led by Drs. Josse and Yang in the 2020 causal inference program. We would like to acknowledge the helpful discussions during the SAMSI working group meetings. We also would like to acknowledge the discussions and insights from the Traumabase group and physicians, in particular Drs François-Xavier AGERON, Tobias GAUSS, and Jean-Denis MOYER. In addition, none of the data analysis part could have been done without the help of Dr. Ian Roberts and the CRASH-3 group, who shared with us the clinical trial data. Part of this work was performed while JJ was a visiting researcher at Google Brain Paris. Finally, we would like to warmly thank Issa Dahabreh for his comments, suggestions of additional references, and insightful discussions.

References

- Ackerman, B., C. Lesko, J. Siddique, R. Susukida, and E. Stuart (2020). Generalizing randomized trial findings to a target population using complex survey population data. *arXiv:2003.07500*. 3
- Athey, S., R. Chetty, and G. Imbens (2020). Combining experimental and observational data to estimate treatment effects on long term outcomes. *arXiv preprint arXiv:2006.09676*. 4.2
- Athey, S., J. Tibshirani, and S. Wager (2019). Generalized random forests. *The Annals of Statistics* 47(2), 1148–1178. 6, 7.1.3
- Bareinboim, E. and J. Pearl (2016). Causal inference and the data-fusion problem. *Proceedings of the National Academy of Sciences* 113, 7345–7352.
- Breiman, L. (2001). Random forests. *Machine learning* 45(1), 5–32. 4.1
- Buchanan, A. L., M. G. Hudgens, S. R. Cole, K. R. Mollan, P. E. Sax, E. S. Daar, A. A. Adimora, J. J. Eron, and M. J. Mugavero (2018). Generalizing evidence from randomized trials using inverse probability of sampling weights. *J. R. Statist. Soc. A*, doi: 10.1111/rssa.12357. 1, 3.5, 3.2.1, D
- Campbell, D. T. (1957). Factors relevant to the validity of experiments in social settings. *Psychological Bulletin* 54(4), 297–312. 1
- Cap, A. P. (2019). Crash-3: a win for patients with traumatic brain injury. *The Lancet* 394(10210), 1687 – 1688.

- Carpenter, J. R. and M. G. Kenward (2007). Missing data in randomised controlled trials: a practical guide. 8
- Chernozhukov, V., D. Chetverikov, M. Demirer, E. Duflo, C. Hansen, and W. K. Newey (2016). Double machine learning for treatment and causal parameters. Technical report, Cemmap Working Paper, Centre for Microdata Methods and Practice.
- Cinelli, C. and J. Pearl (2020). Generalizing experimental results by leveraging knowledge of mechanisms. *European Journal of Epidemiology*. 3, 8, 8
- Cochran, W. G. (1968). The effectiveness of adjustment by subclassification in removing bias in observational studies. *Biometrics* 24, 295–313. 3.2.1
- Cole, S. R. and E. A. Stuart (2010). Generalizing evidence from randomized clinical trials to target populations: The ACTG 320 trial. *American Journal of Epidemiology* 172, 107–115. 1, 1, 3.2.1
- Concato, J., N. Shah, and R. Horwitz (2000). Randomized, controlled trials, observational studies, and the hierarchy of research designs. *The New England journal of medicine* 342, 1887–1892. 1
- Correa, J. D., J. Tian, and E. Bareinboim (2018). Generalized adjustment under confounding and selection biases. In *Thirty-Second AAAI Conference on Artificial Intelligence*. 5.1.3
- CRASH-3 (2019). Effects of tranexamic acid on death, disability, vascular occlusive events and other morbidities in patients with acute traumatic brain injury (CRASH-3): a randomised, placebo-controlled trial. *The Lancet* 394(10210), 1713–1723. 7, 7.2, 7.3.1, 8
- Cro, S., T. P. Morris, B. C. Kahan, V. R. Cornelius, and J. R. Carpenter (2020). A four-step strategy for handling missing outcome data in randomised trials affected by a pandemic. 8
- Dahabreh, I. J., S. J. Haneuse, J. M. Robins, S. E. Robertson, A. L. Buchanan, E. A. Stuart, and M. A. Hernán (2019). Study designs for extending causal inferences from a randomized trial to a target population. *arXiv preprint arXiv:1905.07764*.
- Dahabreh, I. J. and M. A. Hernán (2019). Extending inferences from a randomized trial to a target population. *European Journal of Epidemiology* 34(8), 719–722.
- Dahabreh, I. J., S. E. Robertson, and M. A. Hernán (2019). On the relation between g-formula and inverse probability weighting estimators for generalizing trial results. *Epidemiology* 30(6), 807–812. 3.2.2
- Dahabreh, I. J., S. E. Robertson, J. A. Steingrimsson, E. A. Stuart, and M. A. Hernán (2020). Extending inferences from a randomized trial to a new target population. *Statistics in Medicine* 39(14), 1999–2014. 3.2.4, 3
- Dahabreh, I. J., S. E. Robertson, E. J. Tchetgen, E. A. Stuart, and M. A. Hernán (2019a). Generalizing causal inferences from individuals in randomized trials to all trial-eligible individuals. *Biometrics* 75, 685–694. 1, 5
- Dahabreh, I. J., S. E. Robertson, E. J. T. Tchetgen, E. A. Stuart, and M. A. Hernán (2019b). Generalizing causal inferences from individuals in randomized trials to all trial-eligible individuals. *Biometrics* 75, 685–694. 1, 3.4, 3.2.1, D.3.1, D.3.2

- Dahabreh, I. J., J. M. Robins, and M. A. Hernán (2020). Benchmarking observational methods by comparing randomized trials and their emulations. *Epidemiology* 31(5), 614–619. 5
- Deaton, A. and N. Cartwright (2018). Understanding and misunderstanding randomized controlled trials. *Social Science & Medicine* 210, 2–21. 1
- Deaton, A., S. C. Case, N. Côté, J. Drèze, W. Easterly, R. Khera, L. Pritchett, and C. R. Reddy (2019). Randomization in the tropics revisited: a theme and eleven variations. *Randomized controlled trials in the field of development: A critical perspective*. Oxford University Press. Forthcoming. 1
- Dempster, A. P., N. M. Laird, and D. B. Rubin (1977). Maximum likelihood from incomplete data via the em algorithm. *J. R. Stat. Soc. Ser. B.*, 1–38. 7.3.2
- Dewan, Y., E. Komolafe, J. Mejia-Mantilla, P. Perel, I. Roberts, and H. Shakur-Still (2012, 06). CRASH-3: Tranexamic acid for the treatment of significant traumatic brain injury: study protocol for an international randomized, double-blind, placebo-controlled trial. *Trials* 13, 87. 7, 7.2, 7.3.1
- Dong, L., S. Yang, X. Wang, D. Zeng, and J. Cai (2020). Integrative analysis of randomized clinical trials with real world evidence studies. *arXiv preprint arXiv:2003.01242*. 3.2.3, 3.2.4, 3, 6.3, 6.3
- Frieden, T. (2017, 08). Evidence for health decision making - beyond randomized, controlled trials. *New England Journal of Medicine* 377, 465–475. 1, 1
- Green, L. and R. Glasgow (2006, 04). Evaluating the relevance, generalization, and applicability of research issues in external validation and translation methodology. *Evaluation & the health professions* 29, 126–53. 1
- Guo, F. R. and E. Perković (2020). Efficient least squares for estimating total effects under linearity and causal sufficiency. *arXiv preprint arXiv:2008.03481*. 8
- Guo, R., L. Cheng, J. Li, P. R. Hahn, and H. Liu (2018). A survey of learning causality with data: Problems and methods. *arXiv preprint arXiv:1809.09337*. 6.1
- Hahn, P. R., J. S. Murray, C. M. Carvalho, et al. (2020). Bayesian regression tree models for causal inference: regularization, confounding, and heterogeneous effects. *Bayesian Analysis*. 4.1
- Hainmueller, J. (2012). Entropy balancing for causal effects: A multivariate reweighting method to produce balanced samples in observational studies. *Political Analysis* 20, 25–46. 3.2.3
- Hernán, M., B. Sauer, S. Hernández-Díaz, R. Platt, and I. Shrier (2016, 05). Specifying a target trial prevents immortal time bias and other self-inflicted injuries in observational analyses. *Journal of Clinical Epidemiology* 79. 1
- Hernán, M. A. (2018). The c-word: Scientific euphemisms do not improve causal inference from observational data. *American Journal of Public Health* 108(5), 616–619. PMID: 29565659. 1
- Hernán, M. A. and T. J. VanderWeele (2011). Compound treatments and transportability of causal inference. *Epidemiology* 22, 368–77. 1

- Hill, J. L. (2011). Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics* 20(1), 217–240. 4.1
- Hünermund, P. and E. Bareinboim (2019). Causal inference and data-fusion in econometrics. *arXiv preprint arXiv:1912.09104*. 5, 5.1
- Imbens, G. (2019). Potential outcome and directed acyclic graph approaches to causality: Relevance for empirical practice in economics. Technical report, National Bureau of Economic Research. 1
- Imbens, G. W. and D. B. Rubin (2015). *Causal Inference in Statistics, Social, and Biomedical Sciences*. Cambridge UK: Cambridge University Press. 1
- Jiang, W., J. Josse, M. Lavielle, and T. Group (2020). Logistic regression with missing covariates—parameter estimation, model selection and prediction within a joint-modeling framework. *Computational Statistics & Data Analysis* 145, 106907. 7, 7.3.2
- Josse, J., N. Prost, E. Scornet, and G. Varoquaux (2019). On the consistency of supervised learning with missing values. *arXiv preprint arXiv:1902.06931*. 7.1.1, 7.3.2
- Kallus, N., X. Mao, and M. Udell (2018). Causal inference with noisy and missing covariates via matrix factorization. In *Advances in neural information processing systems*, pp. 6921–6932. 7.1.1
- Kallus, N., A. M. Puli, and U. Shalit (2018). Removing hidden confounding by experimental grounding. In *Advances in Neural Information Processing Systems*, pp. 10888–10897. 4.2, 3
- Kang, J. D. and J. L. Schafer (2007). Demystifying double robustness: A comparison of alternative strategies for estimating a population mean from incomplete data. *Statistical Science* 22, 523–539. 3.2.1
- Keiding, N. and T. A. Louis (2016). Perils and potentials of self-selected entry to epidemiological studies and surveys. *J. R. Statist. Soc. A* 179, 319–376. 1
- Kenward, M. (2013). The handling of missing data in clinical trials. *Clinical Investigation* 3(3), 241–250. 8
- Kern, H. L., E. A. Stuart, J. Hill, and D. P. Green (2016). Assessing methods for generalizing experimental impact estimates to target populations. *Journal of research on educational effectiveness* 9(1), 103–127. 3.2.1, 3.2.1, 3.2.2
- Künzel, S. R., J. S. Sekhon, P. J. Bickel, and B. Yu (2019). Metalearners for estimating heterogeneous treatment effects using machine learning. *Proceedings of the national academy of sciences* 116(10), 4156–4165. 4.1
- Künzel, S. R., S. J. Walter, and J. S. Sekhon (2018). Causaltoolbox—estimator stability for heterogeneous treatment effects. *arXiv preprint arXiv:1811.02833*. 4.1, 6.1
- Lê, S., J. Josse, and F. Husson (2008). FactoMineR: A package for multivariate analysis. *Journal of Statistical Software* 25(1), 1–18. 13, G.2
- Leigh, J., G. Collaborators, Y. Guo, K. Deribe, A. Brazinova, and S. Hostiuc (2018, 11). Global, regional, and national disability-adjusted life years(dalys) for 359 diseases and injuries and healthy life expectancy(hale) for 195 countries and territories, 1990–2017: a systematic analysis for the global burden of disease study 2017. *The Lancet* 392, 1859–1922. 7

- Li, P. and E. A. Stuart (2019). Best (but oft-forgotten) practices: missing data methods in randomized controlled nutrition trials. *The American journal of clinical nutrition* 109(3), 504–508. 8
- Linstone, H. and M. Turoff (1975, 01). *The Delphi Method: Techniques and Applications*, Volume 18. 7.1.2
- Lodi, S., A. Phillips, J. Lundgren, R. Logan, S. Sharma, S. Cole, A. Babiker, M. Law, H. Chu, D. Byrne, A. Horban, J. Sterne, K. Porter, C. Sabin, D. Costagliola, S. Abgrall, M. Gill, G. Touloumi, A. Pacheco, and M. Hernán (2019, 05). Effect estimates in randomized trials and observational studies: Comparing apples with apples. *American Journal of Epidemiology* 188. 1, 8
- Louizos, C., U. Shalit, J. M. Mooij, D. Sontag, R. Zemel, and M. Welling (2017). Causal effect inference with deep latent-variable models. In *Advances in Neural Information Processing Systems*, pp. 6446–6456. 7.1.1
- Lu, Y., D. O. Scharfstein, M. M. Brooks, K. Quach, and E. H. Kennedy (2019). Causal inference for comprehensive cohort studies. *arXiv preprint arXiv:1910.03531*.
- Martel Garcia, F. and L. Wantchekon (2010). Theory, external validity, and experimental inference: Some conjectures. *The ANNALS of the American Academy of Political and Social Science* 628(1), 132–147. 1
- Mayer, I., J. Josse, N. Tierney, and N. Vialaneix (2019). R-miss-tastic: a unified platform for missing values methods and workflows. *arXiv preprint arXiv:1908.04822*. 7.1.1
- Mayer, I., E. Sverdrup, T. Gauss, J.-D. Moyer, S. Wager, and J. Josse (2020). Doubly robust treatment effect estimation with missing attributes. *Ann. Appl. Statist.* 14(3), 1409–1431. 7, 7.1, 7.1.1, 7.1.3, 7.3.3, 8
- National Research Council (2012). The prevention and treatment of missing data in clinical trials. *New England Journal of Medicine* 367, 1355–1360. 8
- Neyman, J. (1923). Sur les applications de la théorie des probabilités aux expériences Agricoles: Essay de principe. English translation of excerpts by Dabrowska, D. and Speed, T. *Statistical Science* 5, 465–472. 1
- Nguyen, T., B. Ackerman, I. Schmid, S. Cole, and E. Stuart (2018, 12). Sensitivity analyses for effect modifiers not observed in the target population when generalizing treatment effects from a randomized controlled trial: Assumptions, models, effect scales, data scenarios, and implementation details. *PLOS ONE* 13, e0208795. 1, 3.2.1, 8
- Nguyen, T. Q., C. Ebnesajjad, S. R. Cole, E. A. Stuart, et al. (2017). Sensitivity analysis for an unobserved moderator in rct-to-target-population generalization of treatment effects. *The Annals of Applied Statistics* 11(1), 225–247. 8
- Nie, X. and S. Wager (2017). Quasi-oracle estimation of heterogeneous treatment effects. *arXiv preprint arXiv:1712.04912*. 4.1

- O'Kelly, M. and B. Ratitch (2014). *Clinical trials with missing data: a guide for practitioners*. John Wiley & Sons. 8
- Olschewski, M. and H. Scheurlen (1985). Comprehensive cohort study: an alternative to randomized consent design in a breast preservation trial. *Methods of Information in Medicine* 24(3), 131–134. 2.2
- O'Muircheartaigh, C. and L. V. Hedges (2014). Generalizing from unrepresentative experiments: a stratified propensity score approach. *J. R. Statist. Soc. C* 63, 195–210. 1, 3.2.1, 3.2.1
- Pearl, J. (1995). Causal diagrams for empirical research. *Biometrika* 82(4), 669–688. 1
- Pearl, J. (2009). *Causality* (2 ed.). Cambridge: Cambridge University Press. 5.1
- Pearl, J. and E. Bareinboim (2011). Transportability of causal and statistical relations: A formal approach. In *Data Mining Workshops (ICDMW), 2011 IEEE 11th International Conference on*, pp. 540–547. IEEE. 1
- Peters, J., D. Janzing, and B. Schölkopf (2017). *Elements of causal inference*. The MIT Press. E
- Peysakhovich, A. and A. Lada (2016). Combining observational and experimental data to find heterogeneous treatment effects. *arXiv preprint arXiv:1611.02385*. 4.2
- Powers, S., J. Qian, K. Jung, A. Schuler, N. H. Shah, T. Hastie, and R. Tibshirani (2018). Some methods for heterogeneous treatment effect estimation in high dimensions. *Statistics in Medicine* 37, 1767–1787.
- Prentice, R. and G. Anderson (2008, 02). The women's health initiative: Lessons learned. *Annual review of public health* 29, 131–50. 1
- R Core Team (2018). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. 7.1.1
- Research, M. (2019). EconML: A Python Package for ML-Based Heterogeneous Treatment Effects Estimation. <https://github.com/microsoft/EconML>. 6.1
- Richardson, T. S. and J. M. Robins (2013). Single world intervention graphs (swigs): A unification of the counterfactual and graphical approaches to causality. *Center for the Statistics and the Social Sciences, University of Washington Series. Working Paper* 128, 2013. 1, 5
- Robins, J. (1986). A new approach to causal inference in mortality studies with a sustained exposure period—application to control of the healthy worker survivor effect. *Mathematical Modelling* 7, 1393–1512. 3.2.2
- Rosenbaum, P. R. and D. B. Rubin (1984). Reducing bias in observational studies using subclassification on the propensity score. *Journal of the American Statistical Association* 79, 516–524. 7.1.1
- Rothman, K. J., J. E. Gallacher, and E. E. Hatch (2013, 08). Why representativeness should be avoided. *International Journal of Epidemiology* 42(4), 1012–1014. 1

- Rothwell, P. M. (2005). External validity of randomised controlled trials: “to whom do the results of this trial apply?”. *The Lancet* 365, 82–93. 1, 1
- Rotnitzky, A. and E. Smucler (2019). Efficient adjustment sets for population average treatment effect estimation in non-parametric causal graphical models. *arXiv preprint arXiv:1912.00306*. 8
- Rubin, D. and M. J. van der Laan (2007). A doubly robust censoring unbiased transformation. *The international journal of biostatistics* 3(1). 4.1
- Rubin, D. B. (1974). Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology* 66, 688–701. 1
- Rubin, D. B. (1976). Inference and missing data. *Biometrika* 63, 581–592. 4.2, 7.1.1
- Rudolph, K. E. and M. J. van der Laan (2017). Robust estimation of encouragement design intervention effects transported across sites. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 79, 1509–1525. 1, 3.2.4
- Saul, B. C. and M. G. Hudgens (2020). The calculus of m-estimation in R with geex. *Journal of Statistical Software* 92(2), 1–15. D.3.1, D.3.2
- Seaman, S. and I. White (2014). Inverse probability weighting with missing predictors of treatment assignment or missingness. *Comm. Statist. Theory Methods* 43, 3499–3515. 7.1.1
- Shadish, W. R., T. D. Cook, D. T. Campbell, et al. (2002). *Experimental and quasi-experimental designs for generalized causal inference/William R. Shadish, Thomas D. Cook, Donald T. Campbell*. Boston: Houghton Mifflin,. 1
- Shakur-Still, H., I. Roberts, R. Bautista, J. Caballero, T. Coats, Y. Dewan, H. El-Sayed, G. Tamar, S. Gupta, J. Herrera, B. Hunt, P. Iribhogbe, M. Izurieta, H. Khamis, E. Komolafe, M. Marreiro, J. Mejía-Mantilla, J. J. Miranda, C. Uribe, and S. Yutthakasemsunt (2009, 11). Effects of tranexamic acid on death, vascular occlusive events, and blood transfusion in trauma patients with significant haemorrhage (CRASH-2): A randomised, placebo-controlled trial. *Lancet* 376, 23–32. 7
- Sharma, A., E. Kiciman, et al. (2019). DoWhy: A Python package for causal inference. <https://github.com/microsoft/dowhy>. 6.1
- Stuart, E. A. (2010). Matching methods for causal inference: A review and a look forward. *Statistical Science* 25, 1–21. 3.2.1
- Stuart, E. A., B. Ackerman, and D. Westreich (2018). Generalizability of randomized trial results to target populations: Design and analysis possibilities. *Research on social work practice* 28(5), 532–537. 1
- Stuart, E. A., C. P. Bradshaw, and P. J. Leaf (2015). Assessing the generalizability of randomized trial results to target populations. *Prevention Science* 16, 475–485. 1
- Stuart, E. A., S. R. Cole, C. P. Bradshaw, and P. J. Leaf (2011). The use of propensity scores to assess the generalizability of results from randomized trials. *J. R. Statist. Soc. A* 174, 369–386. 1, 3.5, 3.2.1, 3.2.1

- Sugiyama, M. and M. Kawanabe (2012). *Machine Learning in Non-stationary Environments: Introduction to Covariate Shift Adaptation*. MIT press. 1
- Textor, J., J. Hardt, and S. Knüppel (2011). Dagitty: a graphical tool for analyzing causal diagrams. *Epidemiology* 22(5), 745. 11
- Tibshirani, J., S. Athey, and S. Wager (2020). *grf: Generalized Random Forests*. R package version 1.2.0. 7.3.2
- Tikka, S., A. Hyttinen, and J. Karvanen (2019). Causal effect identification from multiple incomplete data sources: A general search-based approach. *arXiv preprint arXiv:1902.01073*. 5.1.3, 3, 6.2
- Tikka, S. and J. Karvanen (2017). Identifying causal effects with the R package causaleffect. *Journal of Statistical Software* 76(12), 1–30. 6.1, 3, 6.2
- Tipton, E. (2013). Improving generalizations from experiments using propensity score subclassification: Assumptions, properties, and contexts. *Journal of Educational and Behavioral Statistics* 38, 239–266. 1, 2, 3.2.1, 3.2.1
- Twala, B., M. Jones, and D. J. Hand (2008). Good methods for coping with missing data in decision trees. *Pattern Recognition Letters* 29(7), 950–956. 7.1.1, 7.1.3, 7.3.2
- van Buuren, S. (2018). *Flexible Imputation of Missing Data. Second Edition*. Boca Raton, FL: Chapman and Hall/CRC. 7.1.1, 7.1.3
- Vandenbroucke, J. (2009, 05). The hrt controversy: Observational studies and rcts fall in line. *Lancet* 373, 1233–5. 1
- Wager, S. and S. Athey (2018). Estimation and inference of heterogeneous treatment effects using random forests. *Journal of the American Statistical Association* 113(523), 1228–1242. 4.1, 4.2
- Westreich, D., J. K. Edwards, C. R. Lesko, E. Stuart, and S. R. Cole (2017). Transportability of trial results using inverse odds of sampling weights. *American journal of epidemiology* 186(8), 1010–1014. 1
- Witte, J., L. Henckel, M. H. Maathuis, and V. Didelez (2020). On efficient adjustment in causal graphs. *arXiv preprint arXiv:2002.06825*. 8
- Yang, S., X. Wang, and D. Zeng (2020). Elastic integrative analysis of randomized trial and real-world data for treatment heterogeneity estimation. *arXiv preprint arXiv:2005.10579*. 4.2, 3
- Yang, S., D. Zeng, and X. Wang (2020). Improved inference for heterogeneous treatment effects using real-world data subject to hidden confounding. *arXiv preprint arXiv:2007.12922*. 4.2
- Yao, L., Z. Chu, S. Li, Y. Li, J. Gao, and A. Zhang (2020). A survey on causal inference. *arXiv preprint arXiv:2002.02770*. 6.1
- Zheng, W. and M. J. van der Laan (2011). Cross-validated targeted minimum-loss-based estimation. In *Targeted Learning*, pp. 459–474. Springer. 3.2.4

A Randomized controlled trial

This section recalls assumptions and estimators for average treatment estimation in the case of a single RCT. The assumptions for average treatment effect identifiability in RCTs are the SUTVA assumption and assumptions 3.1 (consistency) and 3.2 (random treatment assignment within the RCT). These assumptions allow the average treatment effect to be identifiable. The most intuitive estimators coming from these assumptions is the difference-in-means estimators:

$$\hat{\tau}_{DM} = \frac{1}{n_1} \sum_{A_i=1} Y_i - \frac{1}{n_0} \sum_{A_i=0} Y_i \quad (\text{S1})$$

With n_1 being the number of individuals in the trial that have been treated and n_0 the number of individuals in the trial who have not been treated ($n_0 + n_1 = n$). This estimator is unbiased and \sqrt{n} -consistent if the sample selection is not moderated by the covariates. If not, it is a biased estimation of the population average treatment effect.

B Estimation of ATE in observational data

Let us recall the IPW estimator in the case of observational data:

$$\hat{\tau}^{\text{IPW}} = \frac{1}{m} \sum_{i=1}^m \left\{ \frac{A_i Y_i}{e(X_i)} - \frac{(1 - A_i) Y_i}{1 - e(X_i)} \right\}, \quad (\text{S2})$$

with $e(x) = P(A = 1 | X = x)$ the propensity score, i.e., the probability to be treated given the covariates. The rationale of IPW is to upweight treated observations with a small propensity score (and the other way around) to balance the two groups, treated and non treated, with respect to their covariates.

C Identification formula

This part focuses on the non-nested design only, as it corresponds to the central design of this review.

Identification by the g-formula or regression formula in the target population

Proof S1

$$\begin{aligned} \mathbb{E}[Y(a)] &= \mathbb{E}[\mathbb{E}[Y(a) | X]] && \text{Law of total expectation} \\ &= \mathbb{E}[\mathbb{E}[Y(a) | X, S = 1]] && \text{Assump. 3.4} \\ &= \mathbb{E}[\mathbb{E}[Y(a) | X, S = 1, A = a]] && \text{Assump. 3.4} \\ &= \mathbb{E}[\mathbb{E}[Y | X, S = 1, A = a]] && \text{Assump. 3.1} \end{aligned}$$

This last quantity can be expressed as a function of the distribution of X in the target population:

$$\mathbb{E}[Y(a)] = \int \mathbb{E}[Y | X = x, S = 1, A = a] df(x),$$

where $f(X)$ denotes the distribution of X in the target population.

Identification by weighting

Proof S2

$$\begin{aligned}
\tau &= \mathbb{E}[\tau(X)] && \text{Law of total expectation} \\
&= \mathbb{E}[\tau_1(X)] && \text{Assump. 3.3} \\
&= \mathbb{E}\left[\frac{f(X)}{f(X | S=1)}\tau_1(X) | S=1\right] && \text{Assump. 3.6.}
\end{aligned}$$

Using Bayes' rule, we note that

$$\frac{f(x)}{f(x | S=1)} = \frac{P(S=1)}{P(S=1 | X=x)} = \frac{P(S=1)}{\pi_S(x)}.$$

In this expression, however, it is important to notice that neither $\pi_S(x)$ nor $P(S=1)$ can be estimated from the data, because we do not observe the S indicator in the observational study (Figure 1). On the other hand, the conditional odds ratio $\alpha(x)$ can be estimated by fitting a logistic regression model that discriminates RCT versus observational samples, and Bayes's rule gives:

$$\begin{aligned}
\alpha(x) &= \frac{P(i \in \mathcal{R} | \exists i \in \mathcal{R} \cup \mathcal{O}, X_i = x)}{P(i \in \mathcal{O} | \exists i \in \mathcal{R} \cup \mathcal{O}, X_i = x)} \\
&= \frac{P(i \in \mathcal{R})}{P(i \in \mathcal{O})} \times \frac{P(X_i = x | i \in \mathcal{R})}{P(X_i = x | i \in \mathcal{O})} \\
&= \frac{n}{m} \times \frac{f(x | S=1)}{f(x)},
\end{aligned}$$

and therefore

$$\tau = \mathbb{E}\left[\frac{n}{m\alpha(x)}\tau_1(X) | S=1\right].$$

This quantity can be further developed, underlying $\tau_1(X)$ identification as presented in proof S3.

Proof S3

$$\begin{aligned}
\tau_1(x) &= \mathbb{E}[Y(1) - Y(0) | X = x, S = 1] \\
&= \mathbb{E}[Y(1) | X = x, S = 1] - \mathbb{E}[Y(0) | X = x, S = 1] \\
&= \frac{\mathbb{E}[A | X = x, S = 1]\mathbb{E}[Y(1) | X = x, S = 1]}{e_1(x)} \\
&\quad - \frac{\mathbb{E}[1 - A | X = x, S = 1]\mathbb{E}[Y(0) | X = x, S = 1]}{1 - e_1(x)} \\
&= \frac{\mathbb{E}[AY(1) | X = x, S = 1]}{e_1(x)} - \frac{\mathbb{E}[(1 - A)Y(0) | X = x, S = 1]}{1 - e_1(x)} && \text{Assump. 3.2} \\
&= \frac{\mathbb{E}[AY | X = x, S = 1]}{e_1(x)} - \frac{\mathbb{E}[(1 - A)Y | X = x, S = 1]}{1 - e_1(x)} && \text{Assump. 3.1} \\
&= \mathbb{E}\left[\frac{A}{e_1(x)}Y - \frac{1 - A}{1 - e_1(x)}Y | X = x, S = 1\right].
\end{aligned}$$

D Nested study design

The nested trial design has different impacts on the estimators expressions previously introduced, and even on the causal quantity of interest. In a nested trial design the randomized trial is embedded in a cohort (e.g. a large cohort - considered as a sample from the target population - in which eligible people are proposed to participate in the trial, but if they refuse they are still included in the cohort study). As a consequence, S is the binary indicator for trial participation, with $S = 1$ for participants and $S = 0$ for non-participants. Therefore the sampling probability of non-randomized individuals is known in nested trial designs (Buchanan et al., 2018; ?). Mathematically it means that the quantity $P(S = 1)$ is identifiable. In addition, two causal quantities can be identified: $\mathbb{E}[Y(1) - Y(0)]$ and $\mathbb{E}[Y(1) - Y(0) | S = 0]$. It is important to note that the second quantity can have a scientific interest in order to better understand heterogeneities within the cohort, and variables that influence the sampling selection and/or the treatment effect on the outcome.

D.1 When observational data have no outcome and treatment information

Main estimators, such as IPSW, g-formula, and doubly-robust estimators are presented for the specific case of nested trial design.

D.1.1 IPSW

In this design the weights in the IPSW estimators are different, because the quantity π_S can be estimated directly from the observed data as the indicator S is observed. This allows the IPSW formula to be closer to the classic IPW expression without the need to use the odds ratio to weight data. The IPSW expression is the following:

$$\hat{\tau}^{\text{IPSW-nested}} = \frac{1}{n} \sum_{i=1}^n \frac{n}{n+m} \frac{A_i Y_i}{\hat{\pi}_S(X_i) e_1(X_i)} - \frac{1}{n} \sum_{i=1}^n \frac{n}{n+m} \frac{(1-A_i) Y_i}{\hat{\pi}_S(X_i) (1-e_1(X_i))}. \quad (\text{S3})$$

The normalized version is the following one:

$$\hat{\tau}^{\text{IPSW-nested norm.}} = \frac{\sum_{i=1}^n (\hat{\pi}_S(X_i) e_1(X_i))^{-1} A_i Y_i}{\sum_{i=1}^n (\hat{\pi}_S(X_i) e_1(X_i))^{-1} A_i} - \frac{\sum_{i=1}^n (\hat{\pi}_S(X_i) (1-e_1(X_i)))^{-1} (1-A_i) Y_i}{\sum_{i=1}^n (\hat{\pi}_S(X_i) (1-e_1(X_i)))^{-1} (1-A_i)}. \quad (\text{S4})$$

Proof S4

$$\begin{aligned} \tau &= \mathbb{E}[\tau(X)] && \text{Law of total expectation} \\ &= \mathbb{E}[\tau_1(X)] && \text{Assump. 3.3} \\ &= \mathbb{E}\left[\frac{f(X)}{f(X | S=1)} \tau_1(X) | S=1\right] && \text{Assump. 3.6} \\ &= \mathbb{E}\left[\frac{P(S=1)}{\pi_S(X)} \tau_1(X) | S=1\right] && \text{Bayes law} \\ &= \mathbb{E}\left[\frac{n}{n+m} \pi_S(X_i)^{-1} \tau_1(X) | S=1\right] && P(S=1) = \frac{n}{n+m} \text{ in the nested design} \end{aligned}$$

Where π_S can be estimated directly using the randomized and the non randomized data. τ_1 is further derived as presented in proof S3.

D.1.2 G-formula

The g-formula formulation in the case of nested trial design depends on the causal quantity of interest. When the target population is the causal quantity of interest, then the identification expression is the same as in the non-nested design. But, because $f \neq f_{\cdot|S=0}$, the estimator's expression is slightly different:

$$\hat{\tau}^{g-nested} = \frac{1}{n+m} \sum_{i=1}^{n+m} (\hat{\mu}_{1,1}(X_i) - \hat{\mu}_{0,1}(X_i)), \quad (\text{S5})$$

In the case where the population of interest is the non-randomized one, the identification of the causal quantity of interest is the following:

$$\mathbb{E}[Y^a | S = 0] = \mathbb{E}[\mathbb{E}[Y | X, S = 1, A = a] | S = 0] = \mathbb{E}[\mu_{1,1}(X) - \mu_{0,1}(X) | S = 0] \quad (\text{S6})$$

The Proof S5 details the calculus. And the estimator is the same as presented in equation 5 as the integration is done on the law $f_{\cdot|S=0}$.

Proof S5

$$\begin{aligned} \mathbb{E}[Y(a)|S=0] &= \mathbb{E}[\mathbb{E}[Y(a) | X]|S=0] && \text{Law of total expectation} \\ &= \mathbb{E}[\mathbb{E}[Y(a) | X, S=1]|S=0] && \text{Assump. 3.4} \\ &= \mathbb{E}[\mathbb{E}[Y(a) | X, S=1, A=a]|S=0] && \text{Assump. 3.4} \\ &= \mathbb{E}[\mathbb{E}[Y | X, S=1, A=a]|S=0] && \text{Assump. 3.1} \end{aligned}$$

This last quantity can be expressed as a function of the distribution of X in the non-randomized population:

$$\mathbb{E}[Y(a)] = \int \mathbb{E}[Y | X = x, S = 1, A = a] f(x|S=0) dx$$

where $f(X|S=0)$ denotes the density function of X in the non-randomized population.

D.1.3 Doubly-robust estimator

Similarly to the doubly-robust estimation in the non-nested case (Section 3.2.4), the g-formula and the IPSW methods can be leveraged into a doubly-robust estimator. The AIPSW expression for the nested case is the following:

$$\begin{aligned} \hat{\tau}^{\text{AIPSW-nested}} &= \frac{1}{n+m} \sum_{i=1}^{n+m} \frac{S_i A_i}{\hat{\pi}_S(X_i) e_1(X_i)} (Y_i - \hat{\mu}_{1,1}(X_i)) \\ &\quad - \frac{1}{n+m} \sum_{i=1}^{n+m} \frac{S_i (1 - A_i)}{\hat{\pi}_S(X_i) (1 - e_1(X_i))} (Y_i - \hat{\mu}_{0,1}(X_i)) \\ &\quad + \frac{1}{m+n} \sum_{i=1}^{m+n} \{\hat{\mu}_{1,1}(X_i) - \hat{\mu}_{0,1}(X_i)\}. \end{aligned}$$

D.2 Combining treatment-effect estimates from both sources of data

Under Assumptions 3.1, 3.2 and 3.5 for the RCT and Assumptions 4.1 and 4.2 for the observational data, separate estimators of the ATEs from the two data sources can be constructed. ? considered the ATEs for the comprehensive cohort studies (CCS) which include participants who would like to be randomized, constituting the RCT, and participants who would like to choose the treatment by their preference, constituting the observational sample. In particular, they considered the ATE over the CCS study population τ_2 and the ATE over the trial population τ_1 . Note that τ_2 is different from τ in our setting because τ_2 is defined with respect to the combined RCT and observational sample; while τ is defined with respect to the observational sample only. In order to construct improved estimators by combining study-specific estimators, they derived the optimal influence functions for τ_1 and τ_2 , which suggest that the efficient estimators of τ_1 and τ_2 can be obtained by

$$\begin{aligned}\hat{\tau}_{1,\text{eff}} &= \frac{1}{n} \sum_{i=1}^{n+m} \left[\frac{\hat{\pi}_S(X_i)A_iY_i}{\hat{e}(X_i)} + \left\{ S_i - \frac{A_i\hat{\pi}_S(X_i)}{\hat{e}(X_i)} \right\} \hat{\mu}_1(X_i) \right. \\ &\quad \left. - \frac{\hat{\pi}_S(X_i)(1-A_i)Y_i}{1-\hat{e}(X_i)} - \left\{ S_i - \frac{(1-A_i)\hat{\pi}_S(X_i)}{1-\hat{e}(X_i)} \right\} \hat{\mu}_0(X_i) \right], \\ \hat{\tau}_{2,\text{eff}} &= \frac{1}{n+m} \sum_{i=n}^{n+m} \frac{A_i\{Y_i - \hat{\mu}_1(X_i)\}}{\hat{e}(X_i)} - \frac{(1-A_i)\{Y_i - \hat{\mu}_0(X_i)\}}{1-\hat{e}(X_i)} + \{\hat{\mu}_1(X_i) - \hat{\mu}_0(X_i)\},\end{aligned}$$

where $\hat{e}_1(X_i)$, $\hat{\mu}_{0,1}(X_i)$, and $\hat{\mu}_{1,1}(X_i)$ for units in the RCT are simplified as $\hat{e}(X_i)$, $\hat{\mu}_0(X_i)$, and $\hat{\mu}_1(X_i)$.

D.3 Softwares: Examples of implementations

This part follows Section 6 and proposes specific examples of implementations for the nested design in the case of IPSW and G-formula.

D.3.1 IPSW

The IPSW estimator can be implemented using the available code from Dahabreh et al. (2019). It requires as input a dataframe (here called `study`) which columns represent treatment, denoted by `A` (binary), the RCT indicator, denoted as `S` (binary), the outcome as `Y` (continuous), and the quantitative covariates. The current available code for 3 quantitative covariates denoted X_1 , X_2 , X_3 is presented below. A first function `generate_weights()` estimates the sampling propensity score and the propensity score as logistic regressions, and compute the according weights to each data point. The variance is estimated with the `geex` library (Saul and Hudgens, 2020) through the `m_estimate` function which computes the empirical sandwich variance estimator.

```

1 # Compute selection score model and propensity score in the trial (logit)
2 weights <- generate_weights(Smod = S~X1+X2+X3, Amod = A~X1+X2+X3, study)
3
4 # Use these scores to compute IPSW
5 IOW1 <- IOW1_est(data = weights$dat)
6
7 # Compute the empirical sandwich variance
8 param_start_IOW1 <- c(coef(weights$Smod) , coef(weights$Amod) ,
```

```

9             m1 = IOW1$IOW1_1, m0 = IOW1$IOW1_0, ate = IOW1$IOW1)
10    IOW1_mest <- m_estimate( estFUN = IOW1_EE, data = study,
11      root_control = setup_root_control(start = param_start_IOW1))
12
13  # Format the output
14  IOW1_ate <- extractEST(geex_output = IOW1_mest,
15    est_name ="ate",
16    param_start = param_start_IOW1)
```

The output is:

```

1 print(IOW1_ate)
2   ate           SE
3   -0.16961   0.02751
```

D.3.2 G-formula

The G-formula can also be implemented in the nested design using the available code from Dahabreh et al. (2019). It takes a similar entry as the IPSW previously presented. The variance is estimated with the `geex` library (Saul and Hudgens, 2020) through the `m_estimate` function which computes the empirical sandwich variance estimator.

```

1 # Linear regression cond. outcome mean as a function of covariates on the RCT
2 # Compute ATE on the observational data
3 OM <- OM_est(data = study)
4
5 # Compute the empirical sandwich variance
6 param_start_OM <- c(coef(OM$OM1mod), coef(OM$OM0mod),
7                         m1=OM$OM_1, m0=OM$OM_0, ate=OM$OM)
8 OM_mest <- m_estimate( estFUN = OM_EE, data = study,
9   root_control = setup_root_control(start = param_start_OM))
10
11 # Format the output
12 OM_ate <- extractEST(geex_output = OM_mest, est_name = "ate",
13 param_start = param_start_OM)
```

The output is:

```

1   ate           SE
2   -0.1934   0.0300
```

E Notations and Assumptions in the Structural Causal Model

Structural Causal Models (SCM). We define a structural causal model as a 3-tuple $M = (U, V, F)$ such that, U and V are sets of respectively unobserved and observed variables with distribution $P(U)$ and $P(V)$, and F is a set of functions $\{f_1, \dots, f_n\}$ such that, for each $v_i \in V$, $v_i = f_i(pa_i, u_i)$, where $pa_i \subseteq V \setminus v_i$ and $u_i \subseteq U$. The notation pa_i (denoted “parents” of v_i) represents the set of variables that directly determine the value of v_i . Such equations are called *structural equations*. For example, let us consider the following SCM M ,

$$a \leftarrow f_a(x, u_a),$$




Figure 15: Left: (a) example of an SCM M and corresponding DAG; right: (b) Post-intervention graph of M for $do(A = a_0)$.

$$y \leftarrow f_y(a, x, u_y).$$

Often, neither parametric assumptions are made on f_a and f_y , nor distributional assumptions on random variables. Each SCM defines a graph G , often a directed acyclic graph (DAG) as illustrated in Figure 15(a).

Interventions. At the core of the SCM framework is the *do-operator* which enables the use of structural equations to represent causal effects and counterfactuals. The $do(A = a_0)$ operation marks the replacements of the mechanism f_a with a constant a_0 , while keeping the rest of the model unchanged, resulting in the following post-treatment model:

$$\begin{aligned} a &\leftarrow a_0 \\ y &\leftarrow f_y(a, x, u_y) \end{aligned}$$

In the causal graph, this corresponds to deleting all incoming arrows in A (Figure 15(b)).

We denote $Q = P(Y = y | do(A = a_0))$ the post-intervention distribution, i.e., the distribution of a random variable Y after a manipulation on A . From this distribution, the ATE can be written as:

$$\begin{aligned} \tau &= \mathbb{E}[Y | do(A = a_1)] - \mathbb{E}[Y | do(A = a_0)] \\ &= \sum_y y(P\{y | do(a_1)\} - P\{y | do(a_0)\}). \end{aligned}$$

Note that the post-intervention distribution can also be denoted in counterfactual notation as $P(y | do(a)) = \overline{P}(Y(a) = y)$. The distinction between $P(Y | A = a)$ and $P(Y | do(a))$ corresponds in the potential outcomes framework to the difference between $P(Y | A = a)$ and $P(Y(a))$.

D-separation. Conditional independences between variables can be read from the DAG induced by an SCM using a graphical criterion known as *d-separation*. This criterion will be useful in identifying the causal effect.

Definition 2 (d-separation) A set X of nodes is said to block a path p if either

- p contains at least one arrow-emitting node that is in X , or

- p contains at least one collision node that is outside X and has no descendant in X .

If X blocks all paths from set A to set Y , it is said to “d-separate A and Y ” and then it can be shown that $A \perp\!\!\!\perp Y | X$. As an illustration, let us consider a path with $A \rightarrow D \leftarrow B \rightarrow C$. Since B emits arrows on that path, it blocks the path between A and C , and $A \perp\!\!\!\perp C | B$. D is a collider (two arrows incoming) and consequently it blocks the path without conditioning $A \perp\!\!\!\perp C$; but conditioning on D would open the path and thus would imply that $A \not\perp\!\!\!\perp C | D$.

Furthermore, in the SCM framework it is generally assumed that *faithfulness* holds, i.e., that all conditional independences are encoded in the graph, allowing to infer dependences from the graph structure (Peters et al., 2017). In other words, if the Global Markov property (i.e., d -separation implies conditional independence), and faithfulness hold, then the resulting equivalence between conditional independences and d -separation allows to move back and forth between the graphical and the probabilistic model.

Identifiability We are interested in answering the *identifiability* question: *can the post-intervention distribution Q be estimated using observed data (such as pre-intervention distribution)?*

Definition 3 (identifiability) A causal query Q is identifiable from distribution $P(y)$ compatible with a causal graph G , if for any two (fully specified) models M_1 and M_2 that satisfy the assumptions in G , we have

$$P_1(v) = P_2(v) \implies Q(M_1) = Q(M_2).$$

Specifically, if a causal query Q in the form of a *do*-expression can be reduced to an expression no longer containing the *do-operator* (i.e, containing only estimable expressions using nonexperimental, observed data) by iteratively applying the inference rules of *do-calculus*, then Q is identifiable. The language of *do-calculus* is proved to be *complete* for queries in the form $Q = P(y | do(a), x)$ meaning that if no reduction can be obtained using these rules, Q is not identifiable.

The application of previous rules and the backdoor criterion in the graph of Figure 16 allows to list all possible admissible adjustment sets for identifying $P(y | do(a))$:

$$\begin{aligned} X = & \{W_2\}, \{W_2, W_3\}, \{W_2, W_4\}, \{W_3, W_4\}, \{W_2, W_3, W_4\}, \{W_2, W_5\}, \{W_2, W_3, W_5\}, \\ & \{W_4, W_5\}, \{W_2, W_4, W_5\}, \{W_3, W_4, W_5\}, \{W_2, W_3, W_4, W_5\} \end{aligned} \quad (S7)$$

The analyst can select from this list which is preferable. Note that conditioning on W_1 would induce bias as it is a collider.

E.1 Proof of the transport formula (Theorem 5.2)

We compute:

$$\begin{aligned} P(y | do(a), S = s^*) &= \sum_x P(y | do(a), S = s^*, X = x) P(x | do(a), S = s^*) \\ &= \sum_x P(y | do(a), X = x, S = s) P(x | do(a), S = s^*) \end{aligned}$$




Figure 16: **Application of the backdoor criterion in large graphs.** Based on the admissible set definition 1, (S7) present all the following sets that are admissible and can be used for adjustment. For example, the set $\{W_2, W_3\}$ blocks all backdoor paths between A and Y . W_2 block the path $A \leftarrow W_2 \rightarrow W_3 \rightarrow W_5 \rightarrow Y$.

$$= \sum_x P(y \mid do(a), X = x, S = s)P(x \mid S = s^*) .$$

The first equation follows by conditioning. The second line is derived by using the s -admissibility of X . The third line results from the fact that A is a child of X and, therefore, exerts no causal influence on X .

F Additional simulation results

This section follows Section 6.3 and provides additional results for the simulations.

F.1 Distributional shift between RCT and observational samples

The simulation design proposed simulates a situation where the RCT data reveals a distributional shift with the observational sample. In the RCT all the covariates tend to have lower values than in the observational sample. Still, the overlap assumption (Assumption 4.2) is still valid as each individual in the target population has a non-zero probability to be included in the experimental sample. Quantitative results obtained for a simulation with ~ 1000 observations in the RCT and 10 000 observations in the observational sample is given on Figure 17, in addition with an histogram illustrating overlaps and the distributional shift for the covariate X_1 .

F.2 Stratification

Within the weighted estimators, the stratification estimator (Section 3.2.1) supposes to choose an additional parameter being the number of strata used. Simulations are launched with the number of strata varying from 3 to 15, and the results are presented on Figure 18. We observed that the number of strata has an impact on the results, the higher the number of strata used, the better the prediction.

	Observational (N=10000)	RCT (N=1023)	Total (N=11023)
X1			
Mean (SD)	1.01 (0.996)	0.552 (0.980)	0.968 (1.00)
Median [Min, Max]	1.01 [-2.84, 4.43]	0.535 [-2.51, 3.62]	0.972 [-2.84, 4.43]
X2			
Mean (SD)	1.00 (0.984)	0.652 (0.991)	0.970 (0.990)
Median [Min, Max]	0.996 [-2.48, 5.02]	0.679 [-2.81, 3.49]	0.963 [-2.81, 5.02]
X3			
Mean (SD)	1.00 (1.01)	0.485 (1.02)	0.954 (1.02)
Median [Min, Max]	1.01 [-2.91, 5.05]	0.468 [-2.32, 3.88]	0.961 [-2.91, 5.05]
X4			
Mean (SD)	0.991 (1.00)	0.616 (1.01)	0.956 (1.01)
Median [Min, Max]	0.988 [-2.77, 4.94]	0.615 [-2.14, 4.17]	0.960 [-2.77, 4.94]




Figure 17: **Covariates distributions differences** between experimental sample and observational sample when simulating according to (15) as detailed in Section 6.3 (left), with a focus on the X_1 distributional shift with histograms overlap for the two samples (right).




Figure 18: **Effect of strata number** Estimated ATE obtained while varying the number of strata $L \in \{3, 5, 7, 9, 11, 13, 15\}$ with 100 repetitions each time. All others simulation parameters being the same as the standard case described in 6.3 and in Figure 6.

F.3 Homogeneous treatment effect

It is always interesting to note that in the case of an homogeneous treatment effect the RCT sample contains all the information to estimate the population ATE, in other words τ_1 is a consistent estimator of the ATE. We performed simulation with an homogeneous treatment effect (results are

presented on Figure (19)) such as:

$$Y(a) | X = -100 + X_1 + 13.7X_2 + 13.7X_3 + 13.7X_4 + 27.4a + \epsilon$$




Figure 19: **Homogeneous treatment effect** Estimated ATE with a homogeneous treatment effect $Y(a) | X = -100 + X_1 + 13.7X_2 + 13.7X_3 + 13.7X_4 + 27.4a + \epsilon$. All others simulation parameters being the same as the standard case described in (6.3) and in Figure 6.

G Additional analysis for Traumabase and CRASH-3

This part proposes additional analysis to the data analysis part (Section 7). We first propose additional visualization of the distributional shift between CRASH-3 and the Traumabase, then we present a principal component analysis of the combined database. Propensity scores obtained either with the logistic regression or the forest are analyzed with histograms and scatter plots. A complementary analysis of the results (Figure 14) is proposed while estimating the generalized treatment effect from an imputed Traumabase on Figure 29. Finally, a focus on the different patients strata, based on the severity of the injury, is presented.

G.1 Distributional shift between CRASH-3 and Traumabase

Distributional shift between CRASH-3 and the Traumabase data can be illustrated with histograms. Figures 20 – 24 presents the empirical distribution shift between the Traumabase and CRASH-3 for age, Glasgow score, systolic blood pressure, sex and pupils reactivity (respectively). Differences can be observed, and for example the fact that the CRASH-3 study contains more young patients, while the Traumabase contains more moderate case (corresponding to a high Glasgow score). It is interesting to notice that the overlaps assumption seems to hold in our situation.




Figure 20: **Distributional shift of Age** between the Traumabase and the CRASH-3 studies.




Figure 21: **Distributional shift of the Glasgow score** between the Traumabase and the CRASH-3 studies.




Figure 22: **Distributional shift of the systolic blood pressure** between the Traumabase and the CRASH-3 studies.




Figure 23: **Distributional shift of the sex** between the Traumabase and the CRASH-3 studies.




Figure 24: **Distributional shift of the pupils reactivity** between the Traumabase and the CRASH-3 studies.

G.2 Principal component analysis

A principal component analysis is performed on the combined data set for the Traumabase and the CRASH-3 data using the **FactoMineR** package (Lê et al., 2008), results are presented on Figure 25. As expected the Glasgow coma scale score and the pupils reactivity are related (paralysis of the cranial nerves leading to pupillary anomalies being closely related to the presence of an intracranial lesion, itself linked to the state of consciousness encoded in the Glasgow.). Additionally, the link between age and systolic blood pressure can be explained by the fact that atherosclerosis of the arteries is the source of an increase in blood pressure and is related to age.




Figure 25: Principal Components Analysis (PCA) of the data set combining CRASH-3 and Traumabase data.

G.3 Sampling propensity scores

The sampling propensity scores obtained while performing the generalization from the CRASH-3 patients to the observational data are presented on Figures 26 (logistic regression) and 27 (forest). We observe that extreme coefficient values are obtained, and that the forest `grf` strengthens this trend. We can further investigate the differences in between the two methods to infer the propensity scores noticing that the forest method uses the `NAs` from the Traumabase to learn the propensity scores model. Figure 28 shows that the `NAs` present in the systolic blood pressure covariate are used by the random forest to predict S , leading to more extreme values at the end. This importance of different missing values patterns when combining two data sets are of importance and highlight the need for a better understanding of the impact of missing values in the present framework.




Figure 26: **Sampling propensity scores histogram (glm)** obtained with the `misaem` R package.




Figure 27: Sampling propensity scores histogram (grf) obtained with random forests.




Figure 28: Scatter plot of the two sampling propensity scores obtained with `glm` in x-axis and `grf` in the y-axis. Color is set according to the systolic blood pressure covariate values (while missing values are in grey).

G.4 Additional results with imputed Traumabase

As the Traumabase presents missing values such that the estimators introduced in this paper (such as IPSW, G-formula, ...) have not been used directly but adapted using methods dealing missing values such as the logistic regression in case of missing values or random forests. Another method consist in imputing the Traumabase and then realize the analysis with estimators introduced in the paper without further modifications. Results are presented on Figure 29 with an imputation using `mice`. Similar conclusions on the treatment effect as the general case presented in Section 7 are obtained.




Figure 29: **Juxtaposition of different estimation results for target population corresponding for all patients** with ATE estimators computed on the **imputed** Traumabase (observational data set), on the CRASH-3 trial (RCT), and transported from CRASH-3 to the Traumabase target population. Number of variables used in each context is given in the legend.

G.5 Evidence on other patients strata

The data analysis part only focuses on all the patients from the two studies CRASH-3 and Traumabase. This part proposes a focus on different patients type, based on the severity of the brain trauma (measured either with the Glasgow score or the pupils reactivity).

G.5.1 Traumabase: evidence on different strata

When stratifying along different criteria of severity as in the CRASH-3 study, namely pupil reactivity and the Glasgow Coma Scale as illustrated in Table 7 with Mild/moderate and Severe strata, the two studies provide different evidence: no average treatment effect in any of the strata for the Traumabase, while the CRASH-3 study finds a beneficial effect for mild forms of TBI.

Table 7: **ATE estimations from the Traumabase** for TBI-related 28-day mortality. Red cells conclude on deteriorating effect, white cells conclude on no effect.

	Multiple imputation (MICE)				MIA		Unadjusted ATE $\times 10^2$
	IPW (95% CI) $\times 10^2$		AIPW (95% CI) $\times 10^2$		IPW (95% CI) $\times 10^2$	AIPW (95% CI) $\times 10^2$	
	GLM	GRF	GLM	GRF			
Total (n = 8248)	15 (6.8, 23)	11 (6.0, 16)	3.4 (-9.0, 16)	-0.1 (-4.7, 4.4)	9.3 (4.0, 15)	-0.4 (-5.2, 4.4)	16
Mild/moderate (GCS > 8, n = 5228)	17 (-7.9, 42)	11 (3.3, 18)	15 (-47, 77)	2.1 (-8.5, 13)	6.8 (2.6, 11)	-0.1 (-4.9, 4.7)	8.7
Severe (GCS ≤ 8, n = 2855)	10 (-7.0, 27)	7.7 (-6.6, 22)	2.2 (-14, 18)	-1.3 (-14, 11)	7.1 (-1.0, 15)	-0.3 (-4.6, 4.0)	9.5

G.5.2 CRASH-3: evidence on different strata

The CRASH-3 trial presents a significant treatment effect only on some strata (in particular on less severe injured patients). As the brain-injury gravity has an effect on the outcome—most patients with TBI with a GCS score of 3 (corresponding to a coma or unconsciousness state) and those with bilateral non-reactive pupils have a very poor prognosis regardless of treatment—the treatment effect is likely to be biased towards the null. Therefore the CRASH-3 authors observe the maximal treatment effect and statistical strength on mild to moderate injured patients, which is what we retrieve from the data. This evidence is computed from the data, with a link between the risk ratio (RR) and the average treatment effect (ATE) on Table 8.

Table 8: **Results reproduction for CRASH-3**, with four possible stratifications based on the gravity level of the injury. Results are both presented as risk ratio (in accordance with CRASH-3 (2019)) and as ATE (in accordance with our framework, Section 2.1).

	Relative risk		Average Treatment Effect	
	RR	95% CI	ATE	95% CI
Total (within 3 hours)	0.94	(0.855, 1.02)	-0.12	(-0.28, 0.004)
GCS > 3 or at least 1 pupil reacts	0.90	(0.78, 1.01)	-0.02	(-0.03, 0.0005)
Mild/moderate (GCS > 8)	0.78	(0.59, 0.98)	-0.2	(-0.03, -0.003)
Severe (GCS ≤ 8)	0.99	(0.91, 1.07)	-0.004	(-0.04, 0.03)
Both pupils react	0.87	(0.74, 1.00)	-0.015	(-0.03, -0.001)

G.5.3 Generalizing treatment effect on patient strata

As found by the CRASH-3 study, the group with potential benefit from TXA seems to be mild to moderate TBI patients (Table 2.1), defined as patients with a Glasgow Coma Scale between 9 and 15, while the evidence obtained from the Traumabase has not found a significant treatment effect for this group. However, in this stratum, for the CRASH-3 study, none of the patients has major extracranial bleeding, leading to a constant variable for this group. Conversely, in the Traumabase, in this stratum, only four patients without major extracranial bleeding are treated (while 1867 are

Table 9: **Sample sizes for both studies** and different strata along the Glasgow Coma Scale. #maj.Ex corresponds to the number of patients with a major extracranial bleeding.

	Traumabase				CRASH-3			
	m	#treated	#death	#maj.Ex	n	#treated	#death	#maj.Ex
Total (within 3 hours)	8248	683	1411	5583	9168	4632	1745	5
Mild/moderate ($GCS > 8$)	5456	535	527	3392	5844	3075	600	0
Severe ($GCS \leq 8$)	3083	596	1322	2224	3717	1985	1601	5

not treated with TXA). Since the practitioners are interested in the treatment effect transported on patients with mild to moderate TBI and with major extracranial bleeding, we cannot restrict the target population to those patients without major extracranial bleeding. The current methodology does not allow to satisfy the necessary assumptions for transporting the effect using the presented estimation strategies and defining a clinically relevant target population. Further methodological investigations are required to transport the effect on the stratified subpopulations (see Table 9 for the corresponding sample sizes).

This issue does not apply to the complementary stratum of severe TBI patients (corresponding to a low Glasgow score ($GCS \leq 8$)). We can thus provide the results for this stratum in Figure 30. We observe that on this strata discrepancies between the solely Traumabase estimators and the generalized estimators are presents. The generalization supports either no-effect or a deleterious effect, while the RCT and the observational estimators support the no-effect hypothesis.




Figure 30: **Juxtaposition of different estimation results for target population corresponding to the severe Traumabase patients** with ATE estimators computed on the Traumabase (observational data set), on the CRASH-3 trial (RCT), and transported from CRASH-3 to the Traumabase target population (severe TBI patients). Number of variables used in each context is given in the legend.